



Chapter 3.

Processes 1.



Operating System Concepts (10th Ed.)





- A *process* is a program in execution.
 - A process is the unit of work in an operating system.
 - A process will need certain resources to accomplish its task.
 - CPU time,
 - memory,
 - files,
 - and I/O devices.





- The memory layout of a process is divided into multiple sections:
 - Text section:
 - the executable code
 - Data section:
 - global variables
 - Heap section:
 - memory that is dynamically allocated during program run time
 - Stack section:
 - temporary data storage when invoking functions
 - such as function parameters, return addresses, and local variables





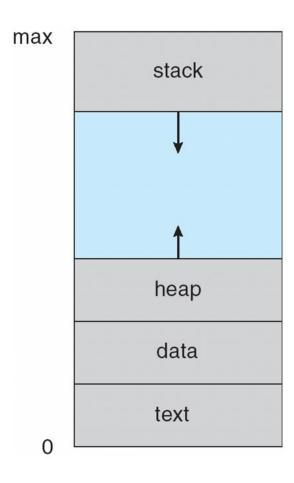
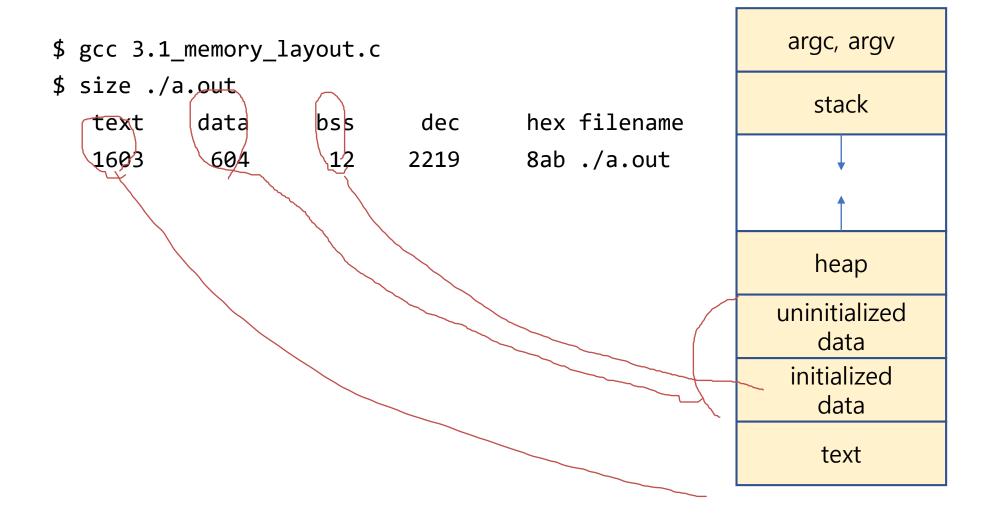


Figure 3.1 Layout of a process in memory.

```
Sowce
#include <stdio.h>`
#include <stdlib.h>
                                                     argc, argv
int x;
                                                        stack
int y = 15;
                                                                       a program in execution
int main(int argc, char *argv[])
    int *values;
                                                       heap
    int i;
                                                    uninitialized
    values = (int *)malloc(sizeof(int)*5);
                                                        data
                                                     initialized
    for (i = 0; i < 5; i++)
                                                        data
        values[i] = i;
                                                        text
    return 0;
```











- As a process executes, it changes its state.
 - New: the process is being created.
 - Running: Instructions are being executed.
 - Waiting: the process is waiting for some event to occur.
 - such as an I/O completion or reception of a signal.
 - Ready: the process is waiting to be assigned to a processor.
 - Terminated: the process has finished execution.



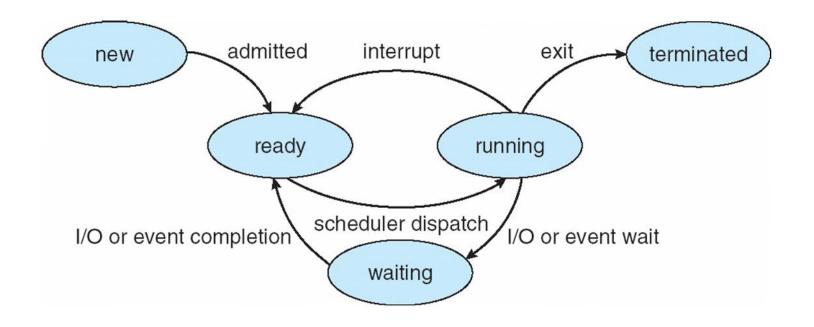


Figure 3.2 Diagram of process state.





- PCB (Process Control Block) or TCB (Task Control Block)
 - Each process is represented in the operating system by the PCB.
- A PCB contains many pieces of information associated with a specific process:
 - Process state
 - Program counter
 - CPU registers
 - CPU-scheduling information
 - Memory-management information
 - Accounting information
 - I/O status information



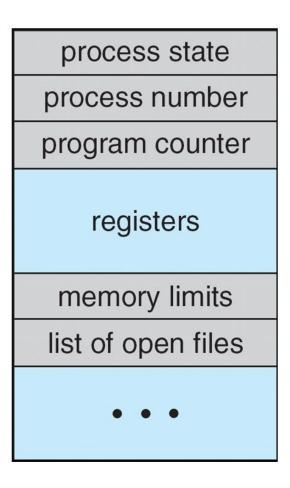


Figure 3.3 *Process control block (PCB).*





- A process is
 - a program that performs a single thread of execution.
 - The single thread of control allows the process to perform
 - only one task at a time.
 - Modern operating systems have extended the process concept
 - to allow a process to have multiple threads of execution
 - and thus to perform more than one task at a time.
- A thread is a lightweight process.
 - Chapter 4 explores multithreading in detail.





- The objective of multiprogramming is
 - to have some process running at all times \geq

```
so as to maximize CPU utilization.
```

- The objective of time sharing is
 - to switch a CPU core among processes so frequently
 - that users can interact with each program while it is running.

```
cpu
```

<-> pararell

at the same time = simultaneously = concurrently



Scheduling Queues:
FIFO: First In First Out

- linkedlist 가
- As processes enter the system, they are put into a ready queue,
 - where they are ready and waiting to execute on a CPU's core.
- Processes that are waiting for a certain event to occur
 - are placed in a wait queue.

 ready queue (running 7) (8
- These queues are generally implemented
 - in the linked lists of PCBs.



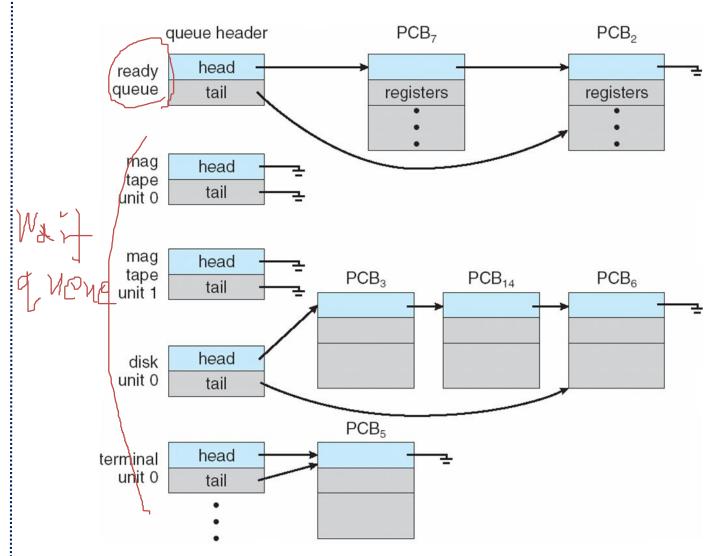


Figure 3.4 *The ready queue and wait queues.*





- Queueing Diagram
 - as a common representation of process scheduling.




Figure 3.5 Queueing-diagram representation of process scheduling.





Context Switch

- The context of a process is represented in the PCB.
- When an interrupt occurs,
 - the system **saves** the current **context** of the running process,
 - so that, later, it can **restore** that **context** when it should be resumed.
- The context switch is a task that
 - switches the CPU core to another process.
 - performs a *state save* of the current process
 - and a *state restore* of a different process.

(PCB)

cpu

<

3.2 Process Scheduling




Figure 3.6 Diagram showing context switch from process to process.





- An operating system must provide a mechanism for
 - process creation,
 - and process termination.
- A process may create several new processes
 - the creating process: a *parent* process.
 - a newly created process: a *child* process.







• A *tree* of processes




Figure 3.7 A tree of processes on a typical Linux system.





- Two possibilities for execution
 - The parent continues to *execute concurrently* with its children.
 - The parent waits until some or all of its children have terminated.
- Two possibilities of address-space
 - The child process is a duplicate of the parent process?
 - The child process has a new program loaded into it.







```
#include <stdio.h> int main()
#include <unistd.h>
#include <wait.h>
                           pid t pid;
                           // fork a child process
                           pid = fork();
                           if (pid < 0) { // error occurred</pre>
                               fprintf(stderr, "Fork Failed");
                               return 1;
                           else if (pid == 0) { // child process
                               execlp("/bin/ls", "ls", NULL);
                           else { // parent process
                               wait(NULL);
                                printf("Child Complete");
                           return 0;
```

Figure 3.8 Creating a separate process using the UNIX fork() system call.





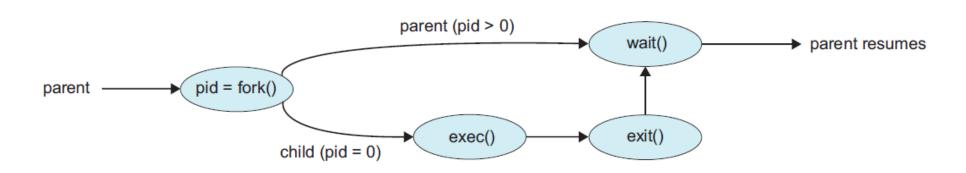


Figure 3.9 Process creation using the fork() system call.





A process terminates.

سل ن

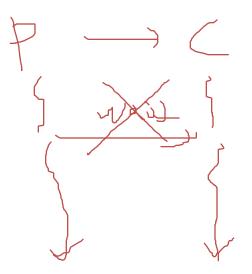



- when it finishes executing its final statement
- exit() system call: asks OS to delete it.
- OS deallocates and reclaims all the resources:
 - allocated memories, open files, and I/O buffers, etc.



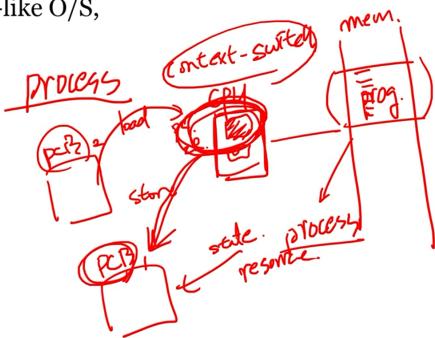
Zombie and Orphan

- zombie process: a process that has terminated,
 - but whose parent has not yet called wait().
- *orphan* process: a process that has a parent process
 - who did not invoke wait() and instead terminated.





■ In UNIX-like O/S,



: process = (RAM) Load

context switch = pcb , pcb load



- In UNIX-like O/S,
 - A new process is created by the fork() system call.
 - The *child* process consists of
 - a **copy** of the address space of the parent process.
 - Both processes continue execution
 - at the instruction after the fork() system call.
 - With one difference:
 - the return code for the fork() is **zero** for the child process, whereas
 - the *nonzero* pid of the child is returned to the parent process.





```
#include <stdio.h>
#include <unistd.h>
int main()
    pid_t pid;
    pid = fork();
    printf("Hello, Process!\n");
```





```
#include <stdio.h>
#include <unistd.h>
int main()
    pid_t pid;
    pid = fork();
    printf("Hello, Process! %d\n", pid);
```





- After a fork() system call,
 - the parent can *continue* its execution; or
 - if it has nothing else to do while the child runs,
 - it can issue a wait() system call
 - to move itself off the ready queue until the termination of the child.



```
#include <stdio.h>
#include <unistd.h>
#include <wait.h>
int main()
    pid_t pid;
    pid = fork();
    if (pid > 0)
        wait(NULL);
    printf("Hello, Process! %d\n", pid);
```



• Exercise 3.1 (p. 154)

```
int value = 5;
int main()
    pid t pid;
    pid = fork();
    if (pid == 0) { // child process
        value += 15;
        return 0;
    else if (pid > 0) { // parent process
        wait(NULL);
        printf("Parent: value = %d\n", value); // LINE A
```

Figure 3.30 What output will be at Line A?





Exercise 3.2 (p. 154)

```
#include <stdio.h>
#include <unistd.h>
#include <wait.h>
/*
 * How many processes are created?
 */
int main()
    fork(); // fork a child process
    fork(); // fork another child process
    fork(); // and fork another
    return 0;
```

Figure 3.31 *How many processes are created?*





• Exercise 3.11 (p. 905)

```
#include <stdio.h>
#include <unistd.h>
/*
 * How many processes are created?
 */
int main()
    int i;
    for (i = 0; i < 4; i++)
        fork();
    return 0;
```

Figure 3.32 *How many processes are created?*





Exercise 3.12 (p. 905)

```
int main()
    pid t pid;
    pid = fork();
    if (pid == 0) { // child process
        execlp("/bin/ls", "ls", NULL);
        printf("LINE J\n");
    else if (pid > 0) { // parent process
        wait(NULL);
        printf("Child Complete\n");
    return 0;
```

Figure 3.33 When will LINE J be reached?





Exercise 3.13 (p. 905)

```
int main()
   pid_t pid, pid1;
   pid = fork();
   if (pid == 0) { // child process
       pid1 = getpid();
       printf("child: pid = %d\n", pid); // A
       printf("child: pid1 = %d\n", pid1); // B
    else if (pid > 0) { // parent process
       pid1 = getpid();
       printf("child: pid = %d\n", pid); // C
       printf("child: pid1 = %d\n", pid1); // D
       wait(NULL);
   return 0;
```

Figure 3.34 What are the pid values?





Exercise 3.16 (p. 905)

```
#define SIZE 5
int nums[SIZE] = \{0, 1, 2, 3, 4\};
int main()
                        if (pid == 0) { // child process
                            for (i = 0; i < SIZE; i++) {
    pid t pid;
    int i;
                                nums[i] *= i;
    pid = fork();
                                printf("CHILD: %d \n", nums[i]); // LINE X
                        else if (pid > 0) { // parent process
    return 0;
                            wait(NULL);
                            for (i = 0; i < SIZE; i++) {
                                printf("PARENT: %d \n", nums[i]); // LINE X
```

Figure 3.35 What output will be at Line X and Line Y?



Any Questions?

