



Chapter 3.

Processes 2.



**Operating System
Concepts (10th Ed.)**





3.4 Interprocess Communication

- Processes executing concurrently may be
 - either *independent* processes or *cooperating* processes.
 - A process is ***independent***
 - if it *does not share data* with any other processes.
 - A process is ***cooperating***
 - if it can *affect* or *be affected* by the other processes.
 - Clearly, any processes that *shares data* with other processes is a cooperating process.



3.4 Interprocess Communication

- **IPC:** *Inter-Process Communication*
 - Cooperating processes require an IPC mechanism
 - that will allow them to exchange data
 - that is, **send data** to and **receive data** from each other.
- Two fundamental models of IPC:
 - *shared memory*
 - *message passing*



3.4 Interprocess Communication

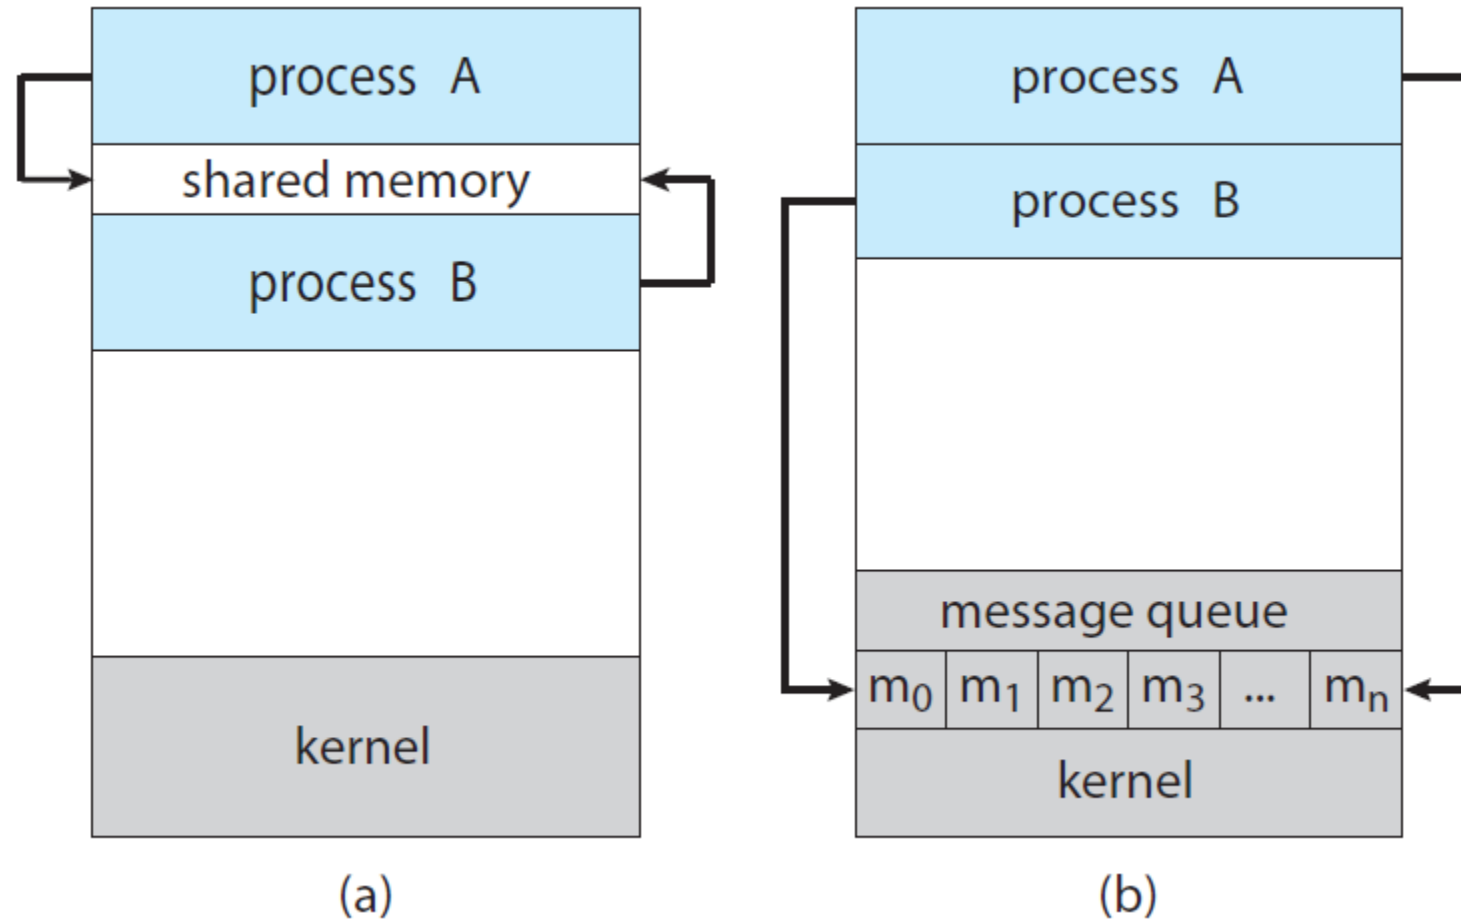


Figure 3.11 Communication models. (a) Shared memory. and (b) Message Passing.

3.5 IPC in Shared-Memory Systems

- Consider the *Producer-Consumer* Problem
 - to illustrate the concept of cooperating processes.
 - a common paradigm for cooperating processes.

- **Producer-Consumer Problem:**
 - A ***producer*** produces information that is consumed by a ***consumer***.
 - For example,
 - a *compiler* produces assembly code, and a *assembler* consumes it.
 - a *web server* produces an HTML file, and a *browser* consumes it.

3.5 IPC in Shared-Memory Systems

- A solution using **shared-memory**:
 - To allow producer and consumer to run *concurrently*.
 - Let a *buffer* of items be available,
 - a producer can *fill the buffer*, and
 - a consumer can *empty the buffer*.
 - A *shared memory* is a region of memory
 - that is shared by the producer and consumer processes.

3.5 IPC in Shared-Memory Systems

- Define a shared buffer
 - in a region of memory shared by the producer and consumer process.

```
#define BUFFER_SIZE 10

typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

3.5 IPC in Shared-Memory Systems

```
item next_produced;

while (true) {
    /* produce an item in next_produced */

    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */

    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}
```

Figure 3.12 *The producer process using shared memory.*

3.5 IPC in Shared-Memory Systems

```
item next_consumed;

while (true) {
    while (in == out)
        ; /* do nothing */

    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next_consumed */
}
```

Figure 3.13 *The consumer process using shared memory.*

3.6 IPC in Message-Passing Systems

- The scheme of using shared-memory
 - requires that these processes *share a region of memory* and that
 - the *code* for accessing and manipulating the shared memory
 - be written *explicitly* by the application programmer.

- **Message-Passing:**
 - O/S provides the means for cooperating processes
 - to communicate with each other via a *message-passing* facility.

3.6 IPC in Message-Passing Systems

- Two operations of the message-passing facility:
 - **send(message)**
 - **receive(message)**

```
message next_produced;

while (true) {
    /* produce an item in next_produced */

    send(next_produced);
}
```

Figure 3.14 *The producer process using message passing.*

3.6 IPC in Message-Passing Systems

```
message next_consumed;

while (true) {
    receive(next_consumed);

    /* consume the item in next_consumed */
}
```

Figure 3.15 *The consumer process using message passing.*

3.6 IPC in Message-Passing Systems

- Communication Links:
 - if two processes P and Q want to communicate,
 - the must *send to* and *receive* messages *from* each other
 - This comm. link can be implemented in a variety of ways.
 - *direct* or *indirect* communication.
 - *synchronous* and *asynchronous* communication.
 - *automatic* or *explicit* buffering.

3.6 IPC in Message-Passing Systems

- Under *direct* communication,
 - each process that wants to communicate
 - must explicitly *name* the *recipient* or *sender* of the communication.
 - The primitives of this scheme:
 - `send(P, message)` – send a message to process *P*.
 - `receive(Q, message)` – receive a message from process *Q*.

3.6 IPC in Message-Passing Systems

- The properties of communication links in this scheme:
 - Links are established *automatically*
 - A link is associated with *exactly two processes*.
 - There exists *exactly one link* between each pair of processes.

3.6 IPC in Message-Passing Systems

- With *indirect* communication,
 - the messages are *sent to* and *received from* **mailboxes**, or **ports**.
 - A **mailbox** (also referred to as **ports**)
 - can be viewed abstractly as an object
 - *into* which messages can be *placed* by processes, and
 - *from* which messages can be *removed*.
 - The primitives of this scheme:
 - **send(*A*, message)** – send a message to mailbox *A*.
 - **receive(*A*, message)** – receive a message from mailbox *A*.

3.6 IPC in Message-Passing Systems

- The properties of communication links in this scheme:
 - Links are established between a pair of processes
 - only if *both members* of the pair have *a shared mailbox*.
 - A link may be associated with *more than two processes*.
 - A number of *different links may exist*, between each pair of processes
 - with each link corresponding to one mailbox.

3.6 IPC in Message-Passing Systems

- OS provides a mechanism that allows a process to do:
 - *Create* a new mailbox.
 - *Send* and *Receive* messages through the mailbox.
 - *Delete* a mailbox.

3.6 IPC in Message-Passing Systems

- Different design options for implementation:
 - *blocking* or *non-blocking*: **synchronous** or **asynchronous**
 - **Blocking send**: the sender is blocked until the message is received.
 - **Non-blocking send**: the sender sends the message and continues.
 - **Blocking receive**: the receiver blocks until a message is available.
 - **Non-blocking receive**: the receiver retrieves either a valid message or a null message.



3.7 Examples of IPC Systems

- Examples of IPC Systems
 - Shared Memory: *POSIX Shared Memory*
 - POSIX: Portable Operating System Interface (for uniX)
 - Message Passing: *Pipes*
 - One of the earliest IPC mechanisms on UNIX systems.



3.7 Examples of IPC Systems

- POSIX shared memory
 - is organized using memory-mapped files,
 - which associate the region of shared memory with a file
 - First, *create* a shared-memory object:
 - `fd = shm_open(name, O_CREAT | ORDWR, 0666);`
 - Configure the *size* of the object in bytes:
 - `ftruncate(fd, 4096);`
 - Finally, establish a *memory-mapped file*:
 - `mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);`



3.7 Examples of IPC Systems

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/mman.h>

int main()
{
    const int SIZE = 4096;          // the size of shared memory
    const char *name = "OS";        // the name of shared memory
    const char *message_0 = "Hello, ";
    const char *message_1 = "Shared Memory!\n";

    int shm_fd;                     // the file descriptor of shared memory
    char *ptr;                      // pointer to shared memory
```



3.7 Examples of IPC Systems

```
/* create the shared memory object */
shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

/* configure the size of the shared memory */
ftruncate(shm_fd, SIZE);

/* map the shared memory object */
ptr = (char *)mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, shm_fd, 0);

/* write to the shared memory */
sprintf(ptr, "%s", message_0);
ptr += strlen(message_0);
sprintf(ptr, "%s", message_1);
ptr += strlen(message_1);

return 0;
}
```

`$gcc 3.16_shm_producer.c -lrt`

Figure 3.16 *Producer process illustrating POSIX shared-memory API.*



3.7 Examples of IPC Systems

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
#include <sys/mman.h>

int main()
{
    const int SIZE = 4096;        // the size of shared memory
    const char *name = "OS";      // the name of shared memory

    int shm_fd;                  // the file descriptor of shared memory
    char *ptr;                   // pointer to shared memory
```




3.7 Examples of IPC Systems

```
/* create the shared memory object */
shm_fd = shm_open(name, O_RDONLY, 0666);

/* map the shared memory object */
ptr = (char *)mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, shm_fd, 0);

/* read from the shared memory object */
printf("%s", (char *)ptr);

/* remove the shared memory */
shm_unlink(name);

return 0;
}

$gcc 3.16_shm_consumer.c -lrt
```

Figure 3.17 Consumer process illustrating POSIX shared-memory API.



3.7 Examples of IPC Systems

- **Pipes** were
 - one of the first IPC mechanisms in early UNIX systems.
 - A **pipe** acts as a *conduit* allowing two processes to communicate.
- Four issues of pipe implementation:
 - Does the pipe allow *unidirectional* or *bidirectional* communication?
 - In the case of two-way comm., is it *half-duplex* or *full-duplex*?
 - Must a *relationship* exist between the communicating process?
 - such as parent-child
 - Can the pipes communicate *over a network*?



3.7 Examples of IPC Systems

- Two common types of pipes:
 - **Ordinary pipes:**
 - *cannot be accessed from outside* the process that created it.
 - Typically, a *parent* process creates a pipe and uses it to communicate with a *child* process that it created.
 - **Named pipes:**
 - can be accessed *without* a *parent-child* relationship



3.7 Examples of IPC Systems

■ Ordinary Pipes

- allow two processes to communicate in producer-consumer fashion.
 - the producer writes to one end of the pipe (*write end*)
 - the consumer reads from the other end (*read end*)
- *unidirectional*: only *one-way* communication is possible.
- *two-way* communication? use two pipes!




Figure 3.20 File descriptors for an ordinary pipe.



3.7 Examples of IPC Systems

- On UNIX systems,
 - ordinary pipes are constructed using the function:
 - `pipe(int fd[])`
 - `fd[0]`: the read end of the pipe
 - `fd[1]`: the write end



3.7 Examples of IPC Systems

```
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>

#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1

int main()
{
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;

    /* create the pipe */
    pipe(fd);
```



3.7 Examples of IPC Systems

```
pid = fork(); // fork a new process

if (pid > 0) { // parent process
    close(fd[READ_END]);
    /* write to the pipe */
    write(fd[WRITE_END], write_msg, strlen(write_msg) + 1);
    close(fd[WRITE_END]);
}
else if (pid == 0) { // child process
    close(fd[WRITE_END]);
    /* read to the pipe */
    read(fd[READ_END], read_msg, BUFFER_SIZE);
    printf("read %s\n", read_msg);
    close(fd[READ_END]);
}

return 0;
}
```

Figure 3.21 Ordinary pipe in UNIX.



3.8 Communication in Client-Server Systems

- Two other strategies in client-server systems
 - *Sockets*
 - are defined as endpoints for communication.
 - *RPCs* (Remote Procedure Calls)
 - abstracts procedure calls between processes on networked systems.



3.8 Communication in Client-Server Systems

- A socket is
 - identified by an *IP address* concatenated with a *port* number.

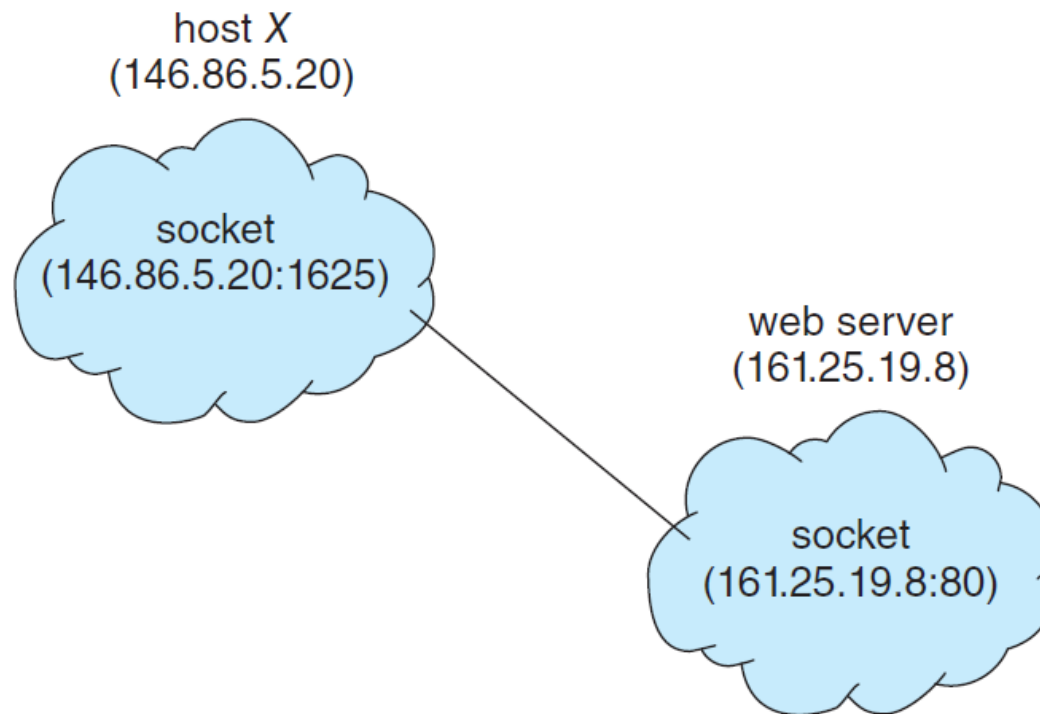


Figure 3.26 Communication using sockets.



3.8 Communication in Client-Server Systems

- Java provides
 - a much easier interface to sockets and
 - provides three different types of sockets
 - `Socket` class: **connection-oriented** (TCP)
 - `DatagramSocket` class: **connectionless** (UDP)
 - `MulticastSocket` class: multiple recipients



3.8 Communication in Client-Server Systems

```
import java.net.*;
import java.io.*;

public class DateServer {

    public static void main(String[] args) throws Exception {
        ServerSocket server = new ServerSocket(6013);

        /* Now listen for connections */
        while (true) {
            Socket client = server.accept();
            PrintWriter pout = new PrintWriter(client.getOutputStream(), true);

            /* write the Date to the socket */
            pout.println(new java.util.Date().toString());

            /* close the socket and resume listening for connections */
            client.close();
        }
    }
}
```

Figure 3.27 *Date server in Java.*



3.8 Communication in Client-Server Systems

```
import java.net.*;
import java.io.*;

public class DateClient {

    public static void main(String[] args) throws Exception {
        /* make connection to server socket */
        Socket socket = new Socket("127.0.0.1", 6013);

        InputStream in = socket.getInputStream();
        BufferedReader br = new BufferedReader(new InputStreamReader(in));

        /* read date from the socket */
        String line = null;
        while ((line = br.readLine()) != null)
            System.out.println(line);

        /* close the socket connections */
        socket.close();
    }
}
```

Figure 3.28 *Date client in Java.*



3.8 Communication in Client-Server Systems

- RPC (Remote Procedure Call)
 - one of the most common forms of *remote service*.
 - designed as a way to abstract the procedure-call mechanism
 - for use between systems with network connections.
 - A client invokes a procedure on a remote host
 - as it would invoke a procedure locally.



3.8 Communication in Client-Server Systems

- The RPC system
 - hides the details that allow communication to take place
 - by providing a *stub* on the client side.
 - The *stub* of client-side locates the server and
 - *marshals* the parameters.
 - The stub of server-side received this message,
 - unpacks the *marshalled* parameters, and
 - performs the procedure on the server.

Any Questions?

