



Chapter 5.

CPU Scheduling



**Operating System
Concepts (10th Ed.)**





5.1 Basic Concepts

- CPU *scheduling* is
 - the basis of multiprogrammed operating systems.
 - The objective of *multiprogramming* is
 - to have some processes running at all times
 - to maximize CPU utilization.



5.1 Basic Concepts


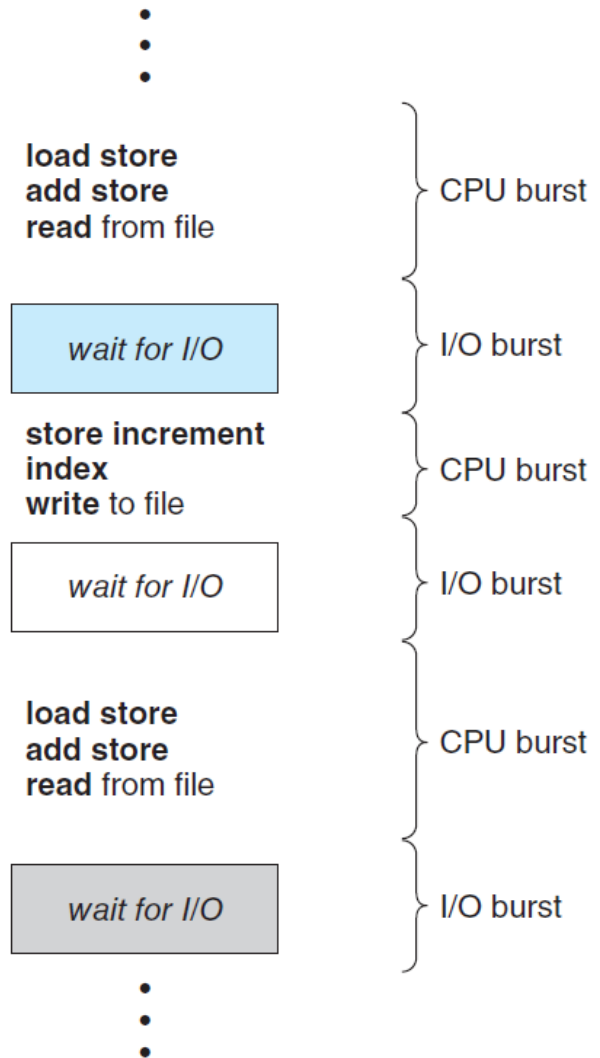


Figure 5.1 *Histogram of CPU-burst duration*

Figure 5.1 *Alternating sequence of CPU-bursts and I/O-bursts.*



5.1 Basic Concepts

- CPU scheduler
 - selects a process from the processes in memory
 - that are *ready* to execute and *allocates* the CPU to that process.
 - Then, how can we select a next process?
 - Linked List? or Binary Tree?
 - *FIFO Queue*: First-In, First-Out
 - *Priority Queue*: How can we determine the priority of a process?



5.1 Basic Concepts

- ***Preemptive*** .vs. ***Non-preemptive***:
 - ***Non-preemptive*** scheduling
 - a process keeps the CPU until it releases it,
 - either by terminating or by switching to the waiting state.
 - ***Preemptive*** scheduling
 - a process can be preempted by the scheduler.



5.1 Basic Concepts

- Decision Making for CPU-scheduling:
 1. When a process switches from the *running* to *waiting* state.
 2. When a process switches from the *running* to *ready* state.
 3. When a process switches from the *waiting* to *ready* state.
 4. When a process *terminates*.
 - No. 1 & 4: no choice – non-preemptive.
 - No. 2 & 3: choices – preemptive or non-preemptive.



5.1 Basic Concepts

- The *dispatcher* is
 - a module that gives control of the CPU's core
 - to the process selected by the CPU scheduler.
 - The functions of dispatcher:
 - switching context from one process to another
 - switching to user mode
 - jumping to the proper location to resume the user program



5.1 Basic Concepts

- The dispatcher should be as fast as possible
 - since it is invoked during every context switch.
 - The *dispatcher latency* is
 - the time to stop one process and start another running.

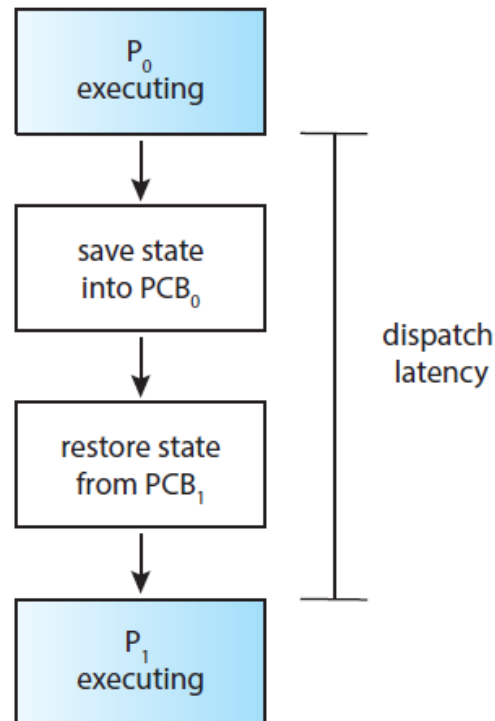


Figure 5.3 *The role of the dispatcher.*



5.1 Basic Concepts

- How *often* do *context switches* occur?

```
joonion@joonionpc:/proc/9$ vmstat 1 3
```

```
procs -----memory-----swap-- ----io---- -system-- -----cpu-----
r  b    swpd      free   buff  cache   si   so    bi    bo    in   cs  us sy id wa st
1  0        0 12879788   7896  53084    0    0     8   791    2    5  0  0 100  0  0
0  0        0 12879796   7896  53084    0    0     0     0    4   30  0  0 100  0  0
0  0        0 12879796   7896  53084    0    0     0     0    5   27  0  0 100  0  0
```

```
joonion@joonionpc:/proc$ cat /proc/1/status | grep ctxt
```

```
voluntary_ctxt_switches:        62
```

```
nonvoluntary_ctxt_switches:    0
```



5.2 Scheduling Criteria

■ Scheduling Criteria

- *CPU utilization*: to keep the CPU as busy as possible.
- *Throughput*: the number of processes completed per time unit.
- *Turnaround time*:
 - how long does it take to execute a process?
 - from the time of submission to the time of completion.
- *Waiting time*:
 - the amount of time that a process spends waiting in the ready queue.
 - the sum of periods spend waiting in the ready queue.
- *Response time*:
 - the time it takes to start responding



5.3 Scheduling Algorithms

- CPU Scheduling Problem:
 - decide which of the processes in the ready queue
 - is to be allocated the CPU's core.



5.3 Scheduling Algorithms

- The solutions for the scheduling problem:
 - **FCFS**: First-Come, First-Served
 - **SJF**: Shortest Job First (**SRTF**: Shortest Remaining Time First)
 - **RR**: Round-Robin
 - **Priority-based**
 - **MLQ**: Multi-Level Queue
 - **MLFQ**: Multi-Level Feedback Queue



5.3 Scheduling Algorithms

■ FCFS Scheduling

- First-Come, First-Served: the simplest CPU-scheduling algorithm.
- The process that requests the CPU first
 - is allocated the CPU first.
 - can be easily implemented with a FIFO queue.



5.3 Scheduling Algorithms


- Consider the following set of processes
 - that arrive at time 0,
 - with the length of the CPU burst given in milliseconds:

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3



5.3 Scheduling Algorithms

- If the processes arrive in the **order** P_1, P_2, P_3 :
 - Gantt Chart served by the FCFS policy:




- Calculate the **waiting time** of this schedule.
 - Waiting Time for $P_1 = 0$, $P_2 = 24$, $P_3 = 27$
 - **Total** Waiting Time: $(0 + 24 + 27) = 51$
 - **Average** Waiting Time: $51/3 = 17$



5.3 Scheduling Algorithms

- If the processes arrive in the **order** P_2, P_3, P_1 :
 - Gantt Chart served by the FCFS policy:




- Calculate the **waiting time** of this schedule.
 - Waiting Time for $P_1 = 6$, $P_2 = 0$, $P_3 = 3$
 - **Total** Waiting Time: $(6 + 0 + 3) = 9$
 - **Average** Waiting Time: $9/3 = 3$




5.3 Scheduling Algorithms

- Let us calculate the *turnaround time*:



- Turnaround Time for $P_1 = 24$, $P_2 = 27$, $P_3 = 30$
- Total* Turnaround Time: $(24 + 27 + 30) = 81$
- Average* Turnaround Time: $81/3 = 27$



- Do it on your own:



5.3 Scheduling Algorithms

- Note that
 - The average waiting time under the *FCFS* policy
 - is generally *not minimal* and *may vary substantially*
 - if the processes' *CPU-burst times* vary greatly.
 - Preemptive or non-preemptive?
 - The FCFS scheduling algorithm is *non-preemptive*.



5.3 Scheduling Algorithms

- Note also that
 - The performance in a *dynamic* situation:
 - What if we have *one CPU-bound* and *many I/O-bound* processes?
 - *Convoy Effect*:
 - all the other processes wait for the one big process to get off the CPU.
 - results in lower CPU and device utilization than might be possible
 - if the shorter processes were allowed to go first.



5.3 Scheduling Algorithms

■ SJF Scheduling


- Shortest-Job-First: *shortest-next-CPU-burst-first* scheduling.
- SJF associates with each process
 - the length of the process's next CPU burst.
- When the CPU is available,
 - assign it to the process that has the smallest next CPU burst.
- If two or more processes are even,
 - break the tie with the FCFS.



5.3 Scheduling Algorithms


<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

- Using the SJF scheduling,
 - Gantt Chart would be:






5.3 Scheduling Algorithms



- Calculate the *waiting time*:
 - Waiting Time for $P_1 = 3$, $P_2 = 16$, $P_3 = 9$, $P_4 = 0$
 - *Total* Waiting Time: $(3 + 16 + 9 + 0) = 28$
 - *Average* Waiting Time: $28/4 = 7$



5.3 Scheduling Algorithms



- Calculate the *turnaround time*:
 - Turnaround Time for $P_1 = 9$, $P_2 = 24$, $P_3 = 16$, $P_4 = 3$
 - *Total* Turnaround Time: $(9 + 24 + 16 + 3) = 52$
 - *Average* Turnaround Time: $52/4 = 13$



5.3 Scheduling Algorithms

- Note that
 - The **SJF** scheduling algorithm is *provably optimal*,
 - it gives the *minimum average waiting time* for a given set of processes.
 - Moving a short process before a long one
 - decreases the waiting time of the short process
 - more than it increases the waiting time of the long process.
 - Consequently, the average waiting time decreases.



5.3 Scheduling Algorithms

- Can you implement the SJF scheduling?
 - There is *no way to know* the length of the *next CPU burst*.
 - Try to *approximate* the SJF scheduling:
 - We may be able to *predict* the length of the next CPU.
 - Pick a process with the shortest *predicted* CPU burst.




5.3 Scheduling Algorithms

- How to predict the next CPU burst?
 - *exponential average* of the measured lengths of *previous* CPU burst.
 - $\tau_{n+1} = \alpha\tau_n + (1 - \alpha)\tau_n$, where
 - τ_n is the length of n th CPU burst,
 - τ_{n+1} is our predicted value for the next CPU burst,
 - for $0 \leq \alpha \leq 1$



5.3 Scheduling Algorithms



CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

$$\alpha = 1/2$$

Figure 5.4 Prediction of the length of the next CPU burst.



5.3 Scheduling Algorithms

- Note also that
 - The SJF algorithm can be either *preemptive* or *non-preemptive*.
 - The choice arises:
 - when a new process arrives at the *ready* queue
 - while a *previous* process is still executing.
 - What if a *newly* arrived process is *shorter* than
 - *what is left* of the currently executing process?



5.3 Scheduling Algorithms


■ SRTF Scheduling

- Shortest-Remaining-Time-First: *Preemptive SJF* scheduling
- SRTF will *preempt* the currently running process,
 - whereas a non-preemptive SJF will *allow* it to *finish* its CPU burst.




5.3 Scheduling Algorithms

- Consider the following:
 - if the processes arrive at the ready queue at the arrival time.





5.3 Scheduling Algorithms



■ The *waiting time*:

- *Total* Waiting Time: $[(10 - 1) + (1 - 1) + (17 - 2) + (5 - 3)] = 26$
- *Average* Waiting Time: $26/4 = 6.5$



5.3 Scheduling Algorithms

- Draw a Gantt chart for non-preemptive SJF scheduling,
 - then calculate the average waiting time to be 7.75.

- Do the same thing with the *turnaround* time.
 - with a *non-preemptive SJF* scheduling algorithm
 - with a *preemptive SRTF* scheduling algorithm.



5.3 Scheduling Algorithms

■ RR Scheduling

- Round-Robin: *preemptive FCFS* with a *time quantum*.
- A *time quantum* (or *time slice*) is a small unit of time.
 - generally from 10 to 100 milliseconds in length.
- The ready queue is treated as a *circular queue*.
- The scheduler goes around the ready queue,
 - allocating the CPU to each process
 - for a time interval of up to 1 time quantum.



5.3 Scheduling Algorithms


- One of two things will happen:
 - The process may have a CPU burst of *less than one time quantum*.
 - the process itself will release the CPU voluntarily
 - the scheduler will proceed to the next process in the ready queue.
 - If the CPU burst is *longer than one time quantum*,
 - the timer will go off and will cause an interrupt to the OS.
 - a context switch will be executed,
 - the process will be put at the tail of the ready queue.



5.3 Scheduling Algorithms


<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- When we use a **time quantum** of 4 milliseconds





5.3 Scheduling Algorithms



■ The *waiting time*:

- Waiting Time for $P_1 = 10 - 4 = 6$, $P_2 = 4$, $P_3 = 7$
- *Total* Waiting Time: $(6 + 4 + 7) = 17$
- *Average* Waiting Time: $17/3 = 5.66$



5.3 Scheduling Algorithms

- Note that
 - The average waiting time under the RR policy is often long.
 - The RR scheduling algorithm is *preemptive*.
 - if a process's CPU burst exceeds one time quantum,
 - that process is *preempted* and is put back in the ready queue.



5.3 Scheduling Algorithms

- The performance of the RR scheduling algorithm
 - depends heavily on the **size** of the *time quantum*.




Figure 5.5 How a smaller time quantum increases context switches.



5.3 Scheduling Algorithms




Figure 5.6 *How turnaround time varies with the time quantum.*



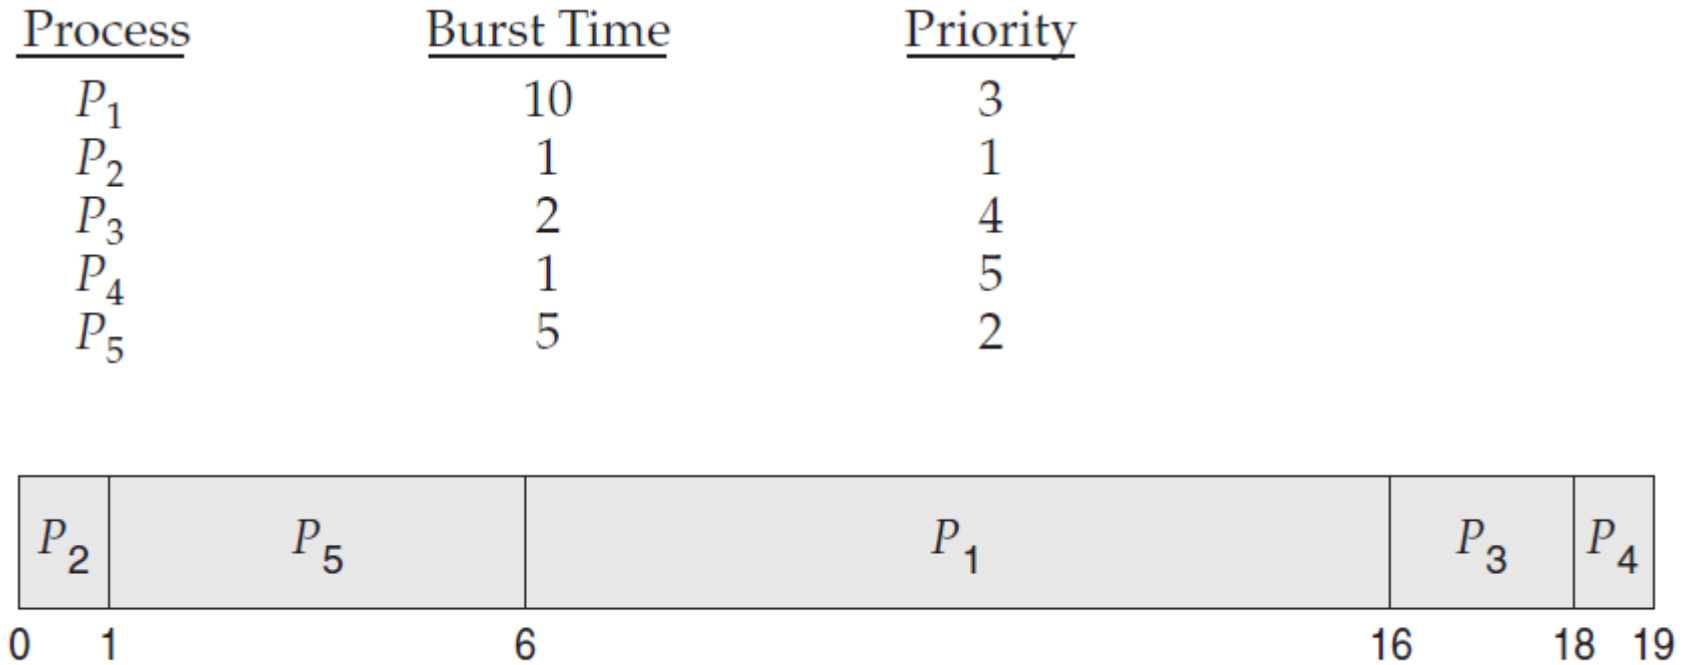
5.3 Scheduling Algorithms

■ Priority-base Scheduling

- A *priority* is associated with each process,
 - and the CPU is allocated to the process with the *highest priority*.
 - Processes with *equal* priority are scheduled in **FCFS** order.
- Note that the ***SJF*** is a *special case* of the *priority-based* scheduling.
 - in this case, the priority is the *inverse* of the *next CPU burst*.
- We assume that *low* numbers represent *high* priority.



5.3 Scheduling Algorithms



- the average waiting time: 8.2
- the average turnaround time: ?



5.3 Scheduling Algorithms

- Priority scheduling can be
 - either *preemptive* or *non-preemptive*.
- The problem of **starvation** (*indefinite blocking*)
 - a *blocked process*: ready to run, but waiting for the CPU.
 - some *low-priority* processes may *wait indefinitely*.
- A solution to the starvation problem is **aging**
 - gradually increase the priority of processes
 - that wait in the system for a long time.




5.3 Scheduling Algorithms

- Combine RR and Priority scheduling:
 - execute the *highest-priority* process and
 - runs processes with the *same* priority using *round-robin* scheduling.

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	4	3
P_2	5	2
P_3	8	2
P_4	7	1
P_5	3	3

- time quantum = 2





5.3 Scheduling Algorithms

■ Multi-Level Queue(MLQ) Scheduling




Figure 5.7 *Separate queues for each priority.*



5.3 Scheduling Algorithms

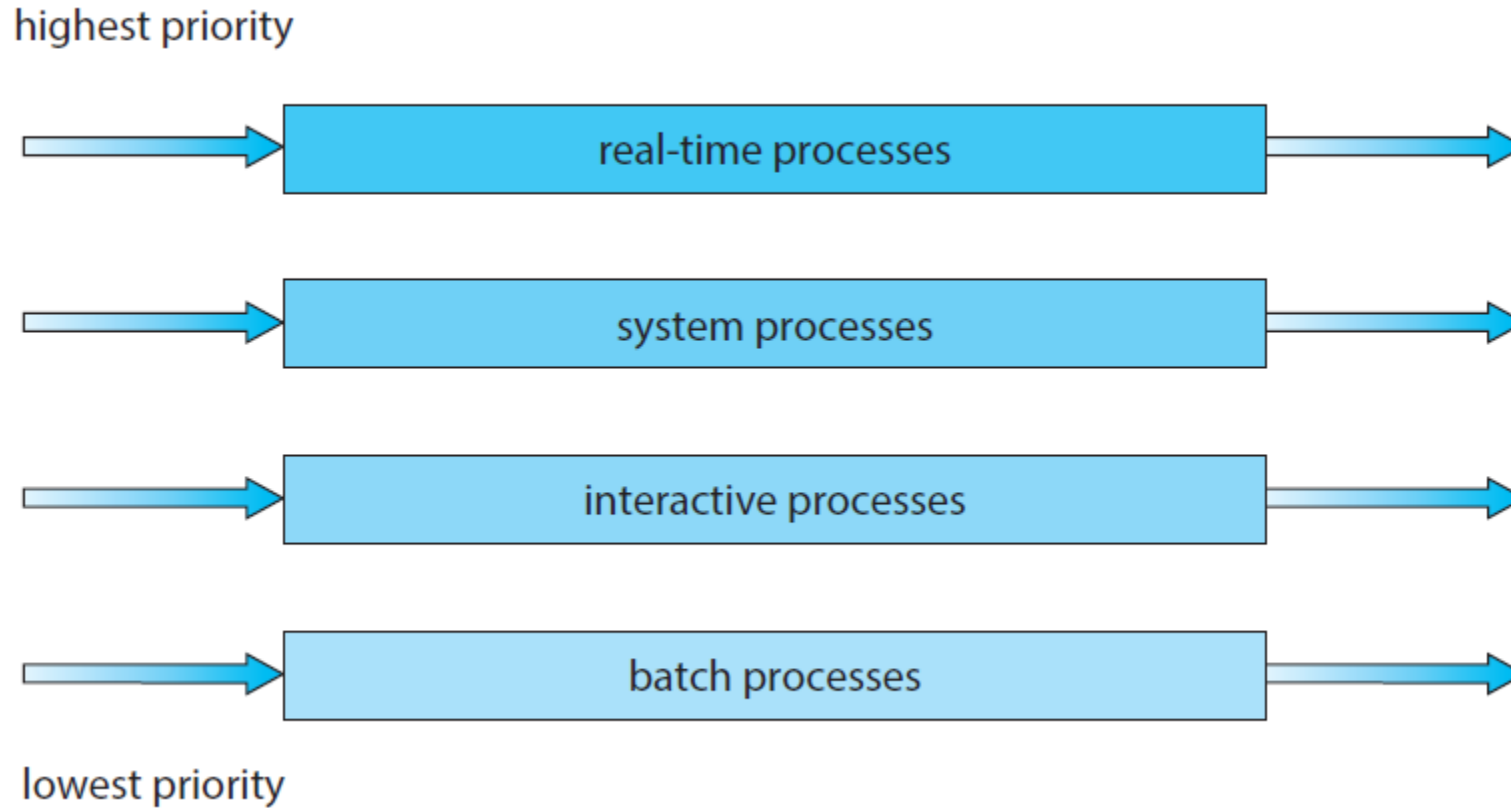


Figure 5.8 *Multi-Level Queue scheduling.*



5.3 Scheduling Algorithms

■ Multi-Level Feedback Queue(MLFQ) Scheduling

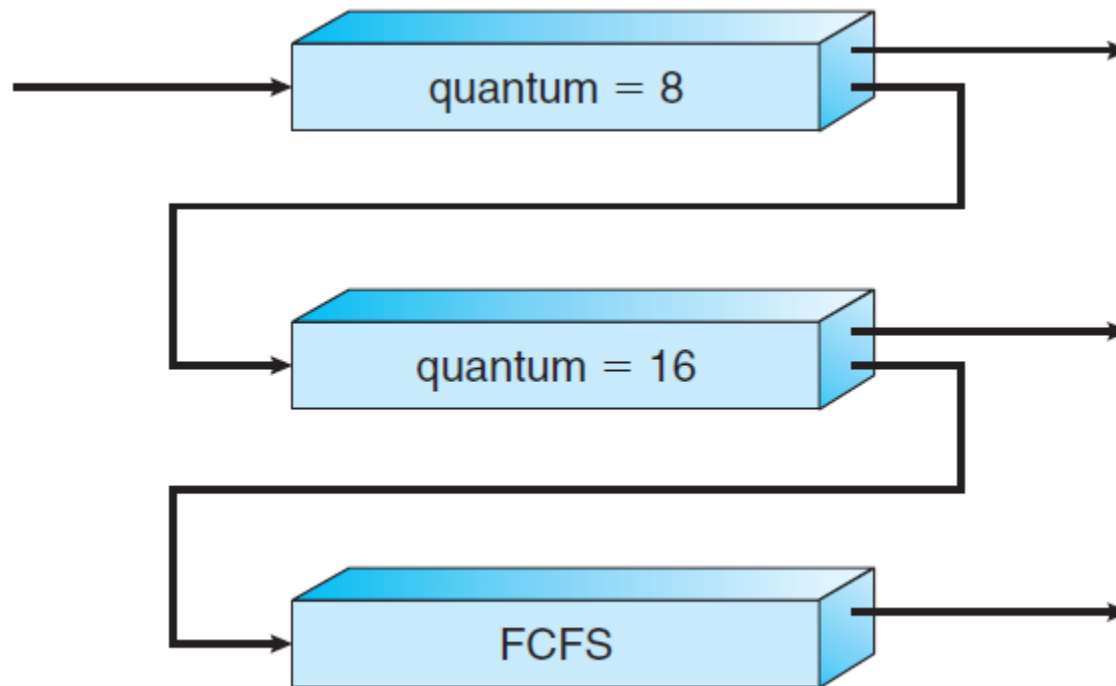


Figure 5.9 *Multi-Level Feedback Queue scheduling.*



5.4 Thread Scheduling

- On most modern operating systems
 - it is *kernel threads* – not *processes* – that are being scheduled,
 - and *user threads* are managed by a thread library.
 - So, the kernel is unaware of them,
 - ultimately mapped to associated kernel threads.



5.6 Real-Time CPU Scheduling

- Scheduling in the Real-Time Operating System.
 - *Soft Realtime* .vs. *Hard Realtime*
 - Soft real-time systems provide no guarantee
 - as to when a critical real-time process will be scheduled.
 - guarantee only that a critical process is preferred to noncritical one.
 - Hard real-time systems have stricter requirements.
 - A task must be services by its deadline.



Practice Exercises

■ Exercise 5.3 (p.251)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	8
P_2	0.4	4
P_3	1.0	1

- What is the average turnaround time for these processes with the FCFS scheduling algorithm?
- What is the average turnaround time for these processes with the SJF scheduling algorithm?



Practice Exercises

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	8
P_2	0.4	4
P_3	1.0	1

- c. The SJF algorithm is supposed to improve performance, but notice that we chose to run process P_1 at time 0 because we did not know that two shorter processes would arrive soon. Compute what the average turnaround time will be if the CPU is left idle for the first 1 unit and then SJF scheduling is used. Remember that processes P_1 and P_2 are waiting during this idle time, so their waiting time may increase. This algorithm could be known as *future-knowledge scheduling*.

Any Questions?

