



## Chapter 6.

# Synchronization Tools

Operating System  
Concepts (10<sup>th</sup> Ed.)





## 6.1 Background

- Cooperating processes
  - can either *affect* or *be affected by* each other.
  - can *share a logical address space* or be allowed to *share data*.
  - However, *concurrent access* to *shared data*
    - may result in ***data inconsistency***.
  - Hence, we need to ensure
    - the *orderly execution* of cooperating processes
    - that share a logical address space to *maintain data consistency*.



## 6.1 Background

- The *integrity of data* shared by several processes (or threads)
  - *Concurrent* execution
    - a process may be *interrupted at any point* in its instruction stream.
    - the processing core may be assigned to another process.
  - *Parallel* execution
    - two or more instruction streams (representing different processes)
    - *execute simultaneously on separate processing cores.*



## 6.1 Background

- Consider an example of how this is happen:
  - Let us revisit the **producer-consumer problem**,
    - where two processes *share data* and are *running asynchronously*.
  - To count items in the buffer, add an integer variable **count**:
    - initialized to 0,
    - *incremented* every time we *add a new item* to the buffer,
    - *decremented* every time we *remove one item* from the buffer.



## 6.1 Background

5

```
while (true) {  
    /* produce an item in next_produced */  
  
    while (count == BUFFER_SIZE)  
        ; /* do nothing */  
  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

```
while (true) {  
    while (count == 0)  
        ; /* do nothing */  
  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
  
    /* consume the item in next_consumed */  
}
```



## 6.1 Background

- Data inconsistency:
  - Although two processes are correct *separately*,
    - they *may not function correctly* when executed *concurrently*.
  - Suppose that the value of count is currently 5,
    - the producer and consumer concurrently execute
    - two statements: `count++;` and `count--;`
  - Then, the value of the variable count may be 4, 5, or 6!
    - is it possible? why?



## 6.1 Background

7

### ■ 다음 프로그램의 출력값은?

```
#include <stdio.h>
#include <pthread.h>

int sum;

void *run(void *param)
{
    int i;
    for (i = 0; i < 10000; i++)
        sum++;
    pthread_exit(0);
}
```

```
int main()
{
    pthread_t tid1, tid2;
    pthread_create(&tid1, NULL, run, NULL);
    pthread_create(&tid2, NULL, run, NULL);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("%d\n", sum);
}
```



## 6.1 Background

8

### ■ 다음 프로그램의 출력값은?

```
int sum;

void *run1(void *param)
{
    int i;
    for (i = 0; i < 10000; i++)
        sum++;
    pthread_exit(0);
}

void *run2(void *param)
{
    int i;
    for (i = 0; i < 10000; i++)
        sum--;
    pthread_exit(0);
}
```

```
int main()
{
    pthread_t tid1, tid2;
    pthread_create(&tid1, NULL, run1, NULL);
    pthread_create(&tid2, NULL, run2, NULL);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("%d\n", sum);
}
```





## 6.1 Background

- How these results can happen?
  - Note that two statements “`count++`” and “`count--`”
    - may be implemented in *machine language* as follows:

```
register1 = count  
register1 = register1 + 1  
count = register1
```

```
register2 = count  
register2 = register2 - 1  
count = register2
```

- Even though *register<sub>1</sub>* and *register<sub>2</sub>* may be the same physical register,
  - the contents of these registers will be
  - ***saved*** and ***restored*** by the *interrupt handler* (or *scheduler*).



## 6.1 Background

- How these results can happen?
  - The concurrent execution of “`count++`” and “`count--`”
    - is equivalent to a sequential execution
    - in which the lower-level statements presented previously
    - are *interleaved* in some *arbitrary order*.

$T_0$ :	<i>producer</i>	execute	$register_1 = count$	$\{register_1 = 5\}$
$T_1$ :	<i>producer</i>	execute	$register_1 = register_1 + 1$	$\{register_1 = 6\}$
$T_2$ :	<i>consumer</i>	execute	$register_2 = count$	$\{register_2 = 5\}$
$T_3$ :	<i>consumer</i>	execute	$register_2 = register_2 - 1$	$\{register_2 = 4\}$
$T_4$ :	<i>producer</i>	execute	$count = register_1$	$\{count = 6\}$
$T_5$ :	<i>consumer</i>	execute	$count = register_2$	$\{count = 4\}$



## 6.1 Background

### ▪ **Race Condition:**

- A situation
  - where *several processes* (or *threads*)
  - access and manipulate the *same* (or *shared*) *data concurrently*
  - and the outcome of the execution
  - *depends on* the *particular order* in which the access takes place.



## 6.1 Background

- To guard against the *race condition*,
  - We need to ensure that
    - *only one process at a time* can manipulate the shared data (e.g. the variable count).
  - To make such a guarantee,
    - we require that the processes are *synchronized* in some way.
    - to say, *process* (or *thread*) *synchronization*.



## 6.1 Background

### ■ Race Condition in Java Threads:

```
public class RaceCondition1 {  
  
    public static void main(String[] args) throws Exception {  
        RunnableOne run1 = new RunnableOne();  
        RunnableOne run2 = new RunnableOne();  
        Thread t1 = new Thread(run1);  
        Thread t2 = new Thread(run2);  
        t1.start(); t2.start();  
        t1.join(); t2.join();  
        System.out.println("Result: " + run1.count + ", " + run2.count);  
    }  
}
```



## 6.1 Background

```
class RunnableOne implements Runnable {  
  
    int count = 0;  
  
    @Override  
    public void run() {  
        for (int i = 0; i < 10000; i++)  
            count++;  
    }  
}
```



## 6.1 Background

```
class RunnableTwo implements Runnable {  
    static int count = 0;  
  
    @Override  
    public void run() {  
        for (int i = 0; i < 10000; i++)  
            count++;  
    }  
}
```



## 6.1 Background

```
public class RaceCondition2 {  
  
    public static void main(String[] args) throws Exception {  
        RunnableTwo run1 = new RunnableTwo();  
        RunnableTwo run2 = new RunnableTwo();  
        Thread t1 = new Thread(run1);  
        Thread t2 = new Thread(run2);  
        t1.start(); t2.start();  
        t1.join(); t2.join();  
        System.out.println("Result: " + RunnableTwo.count);  
    }  
}
```





## Exercises:

- 6.6 Race conditions are possible in many computer systems. Consider a banking system that maintains an account balance with two functions: `deposit(amount)` and `withdraw(amount)`. These two functions are passed the amount that is to be deposited or withdrawn from the bank account balance. Assume that a husband and wife share a bank account. Concurrently, the husband calls the `withdraw()` function, and the wife calls `deposit()`. Describe how a race condition is possible and what might be done to prevent the race condition from occurring.



# Exercises:

- 6.7 The pseudocode of Figure 6.15 illustrates the basic `push()` and `pop()` operations of an array-based stack. Assuming that this algorithm could be used in a concurrent environment, answer the following questions:
- What data have a race condition?
  - How could the race condition be fixed?

```
push(item) {  
    if (top < SIZE) {  
        stack[top] = item;  
        top++;  
    }  
    else  
        ERROR  
}
```

```
pop() {  
    if (!is_empty()) {  
        top--;  
        return stack[top];  
    }  
    else  
        ERROR  
}
```

```
is_empty() {  
    if (top == 0)  
        return true;  
    else  
        return false;  
}
```



## 6.2 The Critical Section Problem

### ■ The Critical Section Problem:

- Consider a system consisting of  $n$  processes  $\{P_0, P_1, \dots, P_{n-1}\}$ .
  - Each process has *a segment of code*, called a **critical section**.
  - in which the process may be *accessing* – and *updating* – *data*
  - that is *shared with* at least one *other process*.
- The important feature of the system is that,
  - when *one process* is *executing* in its *critical section*,
  - *no other process* is *allowed to execute* in its critical section.



## 6.2 The Critical Section Problem

- The **critical-section problem**:
  - No two processes are executing in their critical sections at the same time.
  - To design a protocol that
    - the processes can use to *synchronize* their activity
    - so as to *cooperatively share* data.



## 6.2 The Critical Section Problem

- Sections of codes:
  - The **entry-section**: the section of code
    - to *request permission* to *enter* its critical section.
  - The **critical-section** follows the entry-section.
  - The **exit-section** follows the critical-section.
  - The **remainder-section** is the section of remaining code.



## 6.2 The Critical Section Problem

```
while (true) {  
    entry section  
    critical section  
    exit section  
    remainder section  
}
```

**Figure 6.1** General structure of a typical process.



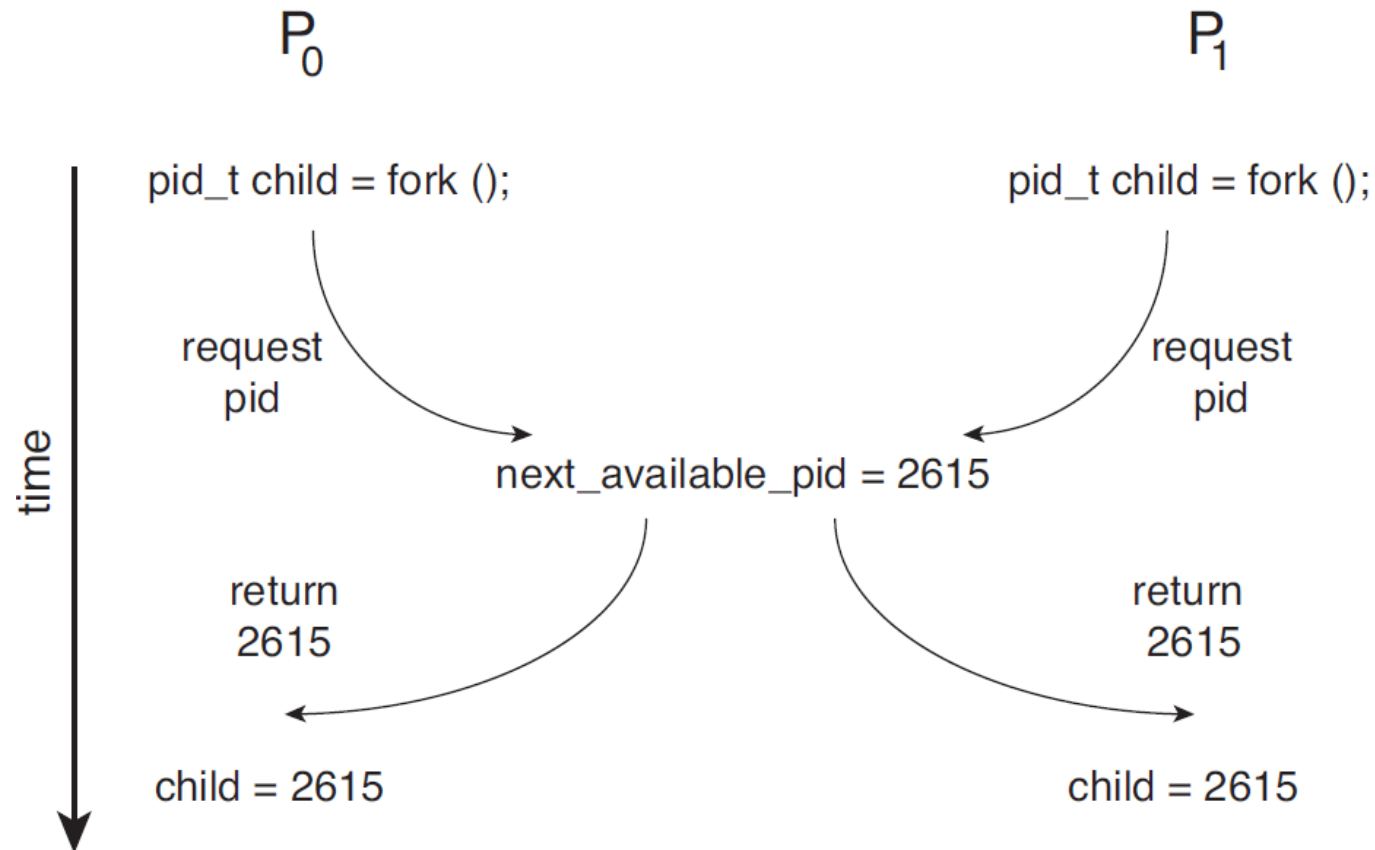
## 6.2 The Critical Section Problem

- Three requirements for the solution:
  - **Mutual Exclusion:**
    - If process  $P_i$  is executing in its critical section,
    - then no other processes can be executing in their critical section.
  - **Progress:** (avoid *deadlock*)
    - If no process is executing in its critical section and some processes wish to enter their critical section,
    - then the selection of next process will enter its critical section next *cannot be postponed indefinitely*.
  - **Bounded Waiting:** (avoid *starvation*)
    - A *bound* (or *limit*) on the number of times that other processes are allowed to enter their critical sections
    - after a process has made a request to enter its critical section and before that request is granted.



## 6.2 The Critical Section Problem

- Example of race condition:



**Figure 6.2** Race condition when assigning a pid.





## 6.2 The Critical Section Problem

- A simple solution in a **single-core** environment:
  - *Prevent interrupts* from occurring
    - while a shared variable was being modified.
  - We could be sure that
    - the current sequence of instructions
    - would be allowed to *execute in order without preemption*.
  - No other instructions would be run,
    - so no unexpected modifications could be made to the shared data.
  - Unfortunately, *not feasible* in a **multiprocessor** environment.



## 6.2 The Critical Section Problem

- Two general approaches:
  - *preemptive* kernels and *non-preemptive* kernels.
  - *Non-preemptive kernel*
    - a kernel-mode process will *run*
    - until it *exits* kernel mode, blocks, or voluntarily yields the CPU.
    - essentially *free from race conditions* on kernel data structures.
  - *Preemptive kernel*
    - *allows* a process to be *preempted* when it is running in kernel mode.
    - essentially *difficult* to design,
    - but *favorable* since it may be more *responsive*.



## 6.3 Peterson's Solution

- Software Solutions to the Critical-Section Problem:
  - *Dekker's* Algorithm:
    - for *two* processes (refer to [Exercise 6.13](#))
  - *Eisenberg and McGuire's* Algorithm:
    - for  $n$  processes with a lower bound on *waiting* of  $n - 1$  turns (refer to [Exercise 6.14](#))
  - **Peterson's** Algorithm:
    - a classic software solution to the critical-section problem.
    - no guarantees that Peterson's solution will work correctly,
    - since modern computers perform basic machine-language instructions
    - such as ***load*** and ***store***.



## 6.3 Peterson's Solution

### ■ Peterson's solution

- restricted to two processes that alternate execution
  - between their *critical* sections and *remainder* sections.

```
while (true) {                                int turn;
    flag[i] = true;                            boolean flag[2];
    turn = j;
    while (flag[j] && turn == j)
        ;

    /* critical section */

    flag[i] = false;

    /*remainder section */
}
```

**Figure 6.3** The structure of process  $P_i$  in Peterson's solution.



## 6.3 Peterson's Solution

- A simple implementation of Peterson's solution:

```
#include <stdio.h>
#include <pthread.h>

#define true 1
#define false 0

int sum = 0;

int turn;
int flag[2];

int main()
{
    pthread_t tid1, tid2;
    pthread_create(&tid1, NULL, producer, NULL);
    pthread_create(&tid2, NULL, consumer, NULL);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("sum = %d\n", sum);
}
```



## 6.3 Peterson's Solution

```
void *producer(void *param)
{
    int k;
    for (k = 0; k < 10000; k++) {
        /* entry section */
        flag[0] = true;
        turn = 1;
        while (flag[1] && turn == 1)
            ;

        /* critical section */
        sum++;

        /* exit section */
        flag[0] = false;

        /* remainder section */
    }
    pthread_exit(0);
}
```

```
void *consumer(void *param)
{
    int k;
    for (k = 0; k < 10000; k++) {
        /* entry section */
        flag[1] = true;
        turn = 0;
        while (flag[0] && turn == 0)
            ;

        /* critical section */
        sum--;

        /* exit section */
        flag[1] = false;

        /* remainder section */
    }
    pthread_exit(0);
}
```



## 6.3 Peterson's Solution

- What happen?
  - There are *no guarantees* that
    - Peterson's solution *will work correctly*,
    - if the architecture perform basic machine-language instructions,
    - such as **load** and **store**.
  - However, Peterson's solution provides
    - a good algorithmic description of solving the CSP.
    - illustrates some of the complexities involved in
    - the requirements of *mutual exclusion*, *progress*, and *bounded waiting*.



## 6.3 Peterson's Solution

- Peterson's solution is *provably correct*.
  - **Mutual exclusion** is preserved.
    - Note that each  $P_i$  enters its critical section,
      - only if either `flag[j]==false` or `turn==i`.
  - The *progress* requirement is satisfied. (**No deadlock**)
  - The *bounded-waiting* requirement is met. (**No starvation**)





## 6.4 Hardware Support for Synchronization

### ■ Hardware-based Solutions

- Hardware instructions that provide
  - support for solving the critical-section problem.
  - can be used *directly* as *synchronization tools*,
  - can be used to form the foundation of *more abstract mechanisms*.
- Three primitive operations
  - ***memory barriers or fences***
  - ***hardware instructions***
  - ***atomic variables***



## 6.4 Hardware Support for Synchronization

### ■ **Atomicity:**

- An *atomic operation* is *one uninterruptible unit* of operation.
- Modern computer systems provide *special hardware instructions*
  - i.e., *atomic instructions*
  - that allow us either to *test and modify* the content of a word
  - or to *test and swap* the contents of two words

### ■ Two types of conceptual atomic instructions:

- `test_and_set()` and `compare_and_swap()`



## 6.4 Hardware Support for Synchronization

- The test\_and\_set() instruction:

---

```
boolean test_and_set(boolean *target) {  
    boolean rv = *target;  
    *target = true;  
  
    return rv;  
}
```

---

**Figure 6.5** *The definition of the atomic test\_and\_set() instruction.*



## 6.4 Hardware Support for Synchronization

- A global Boolean variable `lock`
  - is declared and initialized to `false`.

---

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = false;  
  
    /* remainder section */  
} while (true);
```

---

**Figure 6.6** *Mutual-exclusion implementation with `test_and_set()`.*



## 6.4 Hardware Support for Synchronization

- The `compare_and_swap()` instruction:

---

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
  
    return temp;  
}
```

---

**Figure 6.7** *The definition of the atomic `compare_and_swap()` instruction.*



## 6.4 Hardware Support for Synchronization

- A global Boolean variable `lock`
  - is declared and initialized to 0.

---

```
while (true) {  
    while (compare_and_swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
  
    /* critical section */  
  
    lock = 0;  
  
    /* remainder section */  
}
```

---

**Figure 6.8** *Mutual-exclusion with the `compare_and_swap()` instruction.*



## 6.4 Hardware Support for Synchronization

### ■ Atomic Variable

- Typically, the `compare_and_swap()` instruction
  - is used for construction other tools such as an *atomic variable*.
- An *atomic variable* provides
  - *atomic operations* on *basic data types* such as integers and Booleans.
  - can be used to ensure *mutual exclusion* in situations
  - where there may be a *single variable* with *race condition*.



## 6.4 Hardware Support for Synchronization

- Java implementation of Peterson's solution:

```
public class Peterson1 {  
  
    static int count = 0;  
  
    static int turn = 0;  
    static boolean[] flag = new boolean[2];  
  
    public static void main(String[] args) throws Exception {  
        Thread t1 = new Thread(new Producer());  
        Thread t2 = new Thread(new Consumer());  
        t1.start(); t2.start();  
        t1.join(); t2.join();  
        System.out.println(Peterson1.count);  
    }  
}
```





## 6.4 Hardware Support for Synchronization

```
static class Producer implements Runnable {  
    @Override  
    public void run() {  
        for (int k = 0; k < 10000; k++) {  
            /* entry section */  
            flag[0] = true;  
            turn = 1;  
            while (flag[1] && turn == 1)  
                ;  
  
            /* critical section */  
            count++;  
  
            /* exit section */  
            flag[0] = false;  
  
            /* remainder section */  
        }  
    }  
}
```



## 6.4 Hardware Support for Synchronization

```
static class Consumer implements Runnable {
    @Override
    public void run() {
        for (int k = 0; k < 10000; k++) {
            /* entry section */
            flag[1] = true;
            turn = 0;
            while (flag[0] && turn == 0)
                ;

            /* critical section */
            count--;

            /* exit section */
            flag[1] = false;

            /* remainder section */
        }
    }
}
```



## 6.4 Hardware Support for Synchronization

```
import java.util.concurrent.atomic.AtomicBoolean;

public class Peterson2 {

    static int count = 0;

    static int turn = 0;
    static AtomicBoolean[] flag;
    static {
        flag = new AtomicBoolean[2];
        for (int i = 0; i < flag.length; i++)
            flag[i] = new AtomicBoolean();
    }

}
```



## 6.4 Hardware Support for Synchronization

```
static class Producer implements Runnable {  
    @Override  
    public void run() {  
        for (int k = 0; k < 100000; k++) {  
            /* entry section */  
            flag[0].set(true);  
            turn = 1;  
            while (flag[1].get() && turn == 1)  
                ;  
  
            /* critical section */  
            count++;  
  
            /* exit section */  
            flag[0].set(false);  
  
            /* remainder section */  
        }  
    }  
}
```



## 6.4 Hardware Support for Synchronization

```
static class Consumer implements Runnable {
    @Override
    public void run() {
        for (int k = 0; k < 100000; k++) {
            /* entry section */
            flag[1].set(true);
            turn = 0;
            while (flag[0].get() && turn == 0)
                ;

            /* critical section */
            count--;

            /* exit section */
            flag[1].set(false);

            /* remainder section */
        }
    }
}
```

*Any Questions?*

