



## Chapter 6.

# Synchronization Tools (2)

Operating System  
Concepts (10<sup>th</sup> Ed.)





## 6.5 Mutex Locks

- Higher-level software tools to solve the CSP:
  - *Mutex Locks*: the simplest tools for synchronization.
  - *Semaphore*: more robust, convenient, and effective tool.
  - *Monitor*: overcomes the demerits of mutex and semaphore.
  - *Liveness*: ensures for processes to make progress.



## 6.5 Mutex Locks

- Mutex Lock
  - ***mutex***: ***mut***ual ***ex***clusion.
  - to protect critical section and prevent race condition.
  - a process must ***acquire*** the ***lock*** before *entering* a critical section.
  - ***releases*** the ***lock*** when it *exits* the critical section.



## 6.5 Mutex Locks

- Two functions and one variable for the Mutex Locks:
  - `acquire()` and `release()`
  - `available`: a Boolean variable whose value indicates
    - if the lock is available or not.

```
while (true) {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
}
```

**Figure 6.10** *Solution to the critical-section problem using mutex locks.*



## 6.5 Mutex Locks

- The definition of `acquire()` and `release()`:

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}
```

```
release() {  
    available = true;  
}
```

- Calls to either `acquire()` and `release()` must be performed *atomically*.
- can be implemented using the *compare\_and\_swap* operation.



## 6.5 Mutex Locks

### ■ **Busy waiting:**

- Any other process trying to enter its critical section
  - must *loop continuously* in the call to acquire().
- Busy waiting is clearly a *problem* in a real multiprogramming system,
  - where a single CPU core is shared among many processes.
  - *wastes CPU cycles* for some other processes to use productively.



## 6.5 Mutex Locks

### ■ Spinlock:

- the type of mutex lock using the method of *busy waiting*.
- the process *spins* while waiting for the lock to become available.
- However, spinlocks do have an *advantage*,
  - in that *no context switch* is required waiting on a lock.
  - a context switch may take considerable time.
- In certain circumstances *on multicore systems*,
  - spinlocks are the *preferable* choice for locking.
  - One thread can *spin on one processing core*
  - while another thread performs its critical section *on another core*



## 6.5 Mutex Locks

8

```
void *counter(void *param)
{
    int k;
    for (k = 0; k < 10000; k++) {
        /* entry section */
        pthread_mutex_lock(&mutex);

        /* critical section */
        sum++;

        /* exit section */
        pthread_mutex_unlock(&mutex);

        /* remainder section */
    }
    pthread_exit(0);
}
```

```
#include <stdio.h>
#include <pthread.h>

int sum = 0; // a shared variable

pthread_mutex_t mutex;

int main()
{
    pthread_t tid1, tid2;
    pthread_mutex_init(&mutex, NULL);
    pthread_create(&tid1, NULL, counter, NULL);
    pthread_create(&tid2, NULL, counter, NULL);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("sum = %d\n", sum);
}
```






## 6.6 Semaphores

### ■ Semaphore

- semaphore: 신호장치. 신호기.





## 6.6 Semaphores

### ■ Defining the Semaphore

- A ***semaphore*** *S* is
  - an integer variable that, apart from initialization,
  - is accessed only through *two standard atomic operations*:
  - ***wait()*** and ***signal()***, or sometimes ***P()*** and ***V()***.
- P() and V() are introduced by Edsger Dijkstra
  - ***Proberen***(to test) and ***Verhogen***(to increment)



## 6.6 Semaphores

- Definition of wait() and signal():

```
wait(S) {                               signal(S) {
    while (S <= 0)                        S++;
    ; // busy wait                       }
    S--;
}
```

- All *modifications* to the integer value of the semaphore
  - in the wait() and signal() operations must be executed *atomically*.



## 6.6 Semaphores

- Binary and Counting Semaphores
  - *Binary* Semaphore
    - range only between 0 and 1: similar to *mutex lock*.
  - *Counting* Semaphore
    - range over an unrestricted domain.
    - can be used to resources with *a finite number of instances*.



## 6.6 Semaphores

- Using the counting semaphore:
  - Initialize a semaphore to *the number of resources available*.
  - When a process uses a resource
    - `wait()` on the semaphore: *decrements the count*.
  - When a process release a resource
    - `signal()` on the semaphore: *increment the count*.
  - When the count goes to 0, all resources are being used.
    - Then, processes that wish to use a resource *will block*
    - *until* the count becomes *greater than 0*.



## 6.6 Semaphores

- Using the semaphore to solve synchronization problem:
  - Consider two processes  $P_1$  and  $P_2$  running concurrently.
    - $P_1$  with a statement  $S_1$ , and  $P_2$  with a statement  $S_2$ .
  - Suppose that  $S_2$  should be *executed only after*  $S_1$  has completed.
    - Let  $P_1$  and  $P_2$  share a *semaphore* **synch**, initialized to 0.

```
 $S_1$ ;  
signal(synch);
```

```
wait(synch);  
 $S_2$ ;
```



## 6.6 Semaphores

- Semaphore Implementation:
  - Semaphores also suffer from the problem of *busy waiting*.
  - To overcome this problem, modify the definition of P() and V().
  - When a process executes the **wait()** operation
    - and finds that the semaphore *is not positive*, it must *wait*.
    - rather than busy waiting, suspend itself and goes to the *waiting queue*.
  - When other process executes the **signal()** operation
    - waiting processes can be *restarted* and placed into the *ready queue*.



## 6.6 Semaphores

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        sleep();  
    }  
}
```

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```





## 6.6 Semaphores

```
void *counter(void *param)
{
    int k;
    for (k = 0; k < 10000; k++) {
        /* entry section */
        sem_wait(&sem);

        /* critical section */
        sum++;

        /* exit section */
        sem_post(&sem);

        /* remainder section */
    }
    pthread_exit(0);
}
```

```
#include <stdio.h>
#include <pthread.h>
#include <semaphore.h>

int sum = 0; // a shared variable

sem_t sem;

int main()
{
    pthread_t tid1, tid2;
    sem_init(&sem, 0, 1);
    pthread_create(&tid1, NULL, counter, NULL);
    pthread_create(&tid2, NULL, counter, NULL);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    printf("sum = %d\n", sum);
}
```



## 6.6 Semaphores

```
int main()
{
    pthread_t tid[5]; int i;
    sem_init(&sem, 0, 5);
    for (i = 0; i < 5; i++)
        pthread_create(&tid[i], NULL, counter, NULL);
    for (i = 0; i < 5; i++)
        pthread_join(tid[i], NULL);
    printf("sum = %d\n", sum);
}
```

## 6.7 Monitors

- The *difficulty* of using semaphores:
  - The semaphore is *convenient* and *effective* for synchronization.
  - However, *timing errors* can happen
    - if particular execution sequences take place.
    - these sequences *do not always occur*,
    - and it is *hard to detect*.

## 6.7 Monitors

- An illustrative example of semaphore's problem
  - All processes share a *binary semaphore* **mutex** initialized to 1.
    - Each process must **wait(mutex)** before entering the CS
    - and **signal(mutex)** afterward.
  - If this sequence is not observed,
    - two processes *may be in* their critical sections *simultaneously*.



## 6.7 Monitors

- Situation 1:
  - Note that the difficulty arises
    - even if a single process is not well behaved.
  - Suppose that a program *interchanges the order*.
    - in which **wait()** and **signal()** on the semaphore **mutex** are executed.

```
signal(mutex);  
...  
critical section  
...  
wait(mutex);
```

## 6.7 Monitors

- Situation 2 & 3:
  - Suppose that a program *replaces **signal()** with **wait()***.

```
wait(mutex);  
...  
critical section  
...  
wait(mutex);
```

- Suppose that a process omits the **wait()**, or the **signal()**, or both of them.



## 6.7 Monitors

- How to deal with these kinds of difficulties?
  - These situations may be caused
    - by an honest programming error or an uncooperating programmer.
  - Various types of errors can be generated easily
    - when programmers use semaphores (or mutex locks) incorrectly.
  - Incorporate simple synchronization tools
    - as high-level language constructs
    - **monitor**: one fundamental high-level synchronization construct.

## 6.7 Monitors

- A ***monitor type*** is
  - an **ADT** that includes a set of *programmer-defined operations*
    - that are provided with mutual exclusion within the ***monitor***.
  - declares the ***variables***
    - whose values define the *state of an instance* of that type.
    - along with the bodies of ***function*** that operate on those *variables*.





## 6.7 Monitors

```
monitor monitor name
{
    /* shared variable declarations */

    function P1 ( . . . ) {
        . . .
    }


    function P2 ( . . . ) {
        . . .
    }

    .
    .
    .
    function Pn ( . . . ) {
        . . .
    }

    initialization_code ( . . . ) {
        . . .
    }
}
```

**Figure 6.11** *Pseudocode syntax of a monitor.*

## 6.7 Monitors



**Figure 6.12** *Schematic view of a monitor.*

## 6.7 Monitors

### ■ **Conditional Variables:**

- The monitor construct is not sufficiently powerful
  - for modeling some synchronization schemes.
- We need to define the **condition** construct
  - to provide additional synchronization mechanisms.

## 6.7 Monitors


- Using conditional variables:
  - One can define one or more variables of type **condition**:
  - The only operations that can be invoked
    - on a **condition variable** are **wait()** and **signal()**.

```
condition x, y;
```

```
x.wait();
```

```
x.signal();
```

## 6.7 Monitors



**Figure 6.13** *Monitor with condition variables.*

## 6.7 Monitors

### ■ Java Monitors

- Java provides a *monitor-like*
  - *concurrency* mechanism for *thread synchronization*.
  - called as ***monitor-lock*** or ***intrinsic-lock***.
- Basic language constructs for Java Synchronization
  - **synchronized** keyword.
  - **wait()** and **notify()** method.

## 6.7 Monitors

### ▪ **synchronized** keyword:

- 임계영역에 해당하는 코드 블록을 선언할 때 사용하는 자바 키워드
- 해당 코드 블록(임계영역)에는 모니터락을 획득해야 진입 가능
- 모니터락을 가진 객체 인스턴스를 지정할 수 있음
- 메소드에 선언하면 메소드 코드 블록 전체가 임계영역으로 지정됨
  - 이 때, 모니터락을 가진 객체 인스턴스는 `this` 객체 인스턴스임

```
synchronized (object) {  
    // critical section  
}
```

```
public synchronized void add() {  
    // critical section  
}
```

## 6.7 Monitors

- **wait()** and **notify()** methods:
  - `java.lang.Object` 클래스에 선언됨: 모든 자바 객체가 가진 메소드임
  - 쓰레드가 어떤 객체의 `wait()` 메소드를 호출하면
    - 해당 객체의 모니터락을 획득하기 위해 대기 상태로 진입함.
  - 쓰레드가 어떤 객체의 `notify()` 메소드를 호출하면
    - 해당 객체 모니터에 대기중인 쓰레드 **하나**를 깨움.
  - `notify()` 대신에 `notifyAll()` 메소드를 호출하면
    - 해당 객체 모니터에 대기중인 쓰레드 **전부**를 깨움.





## 6.7 Monitors

### ■ Java Synchronization Example 1:

```
public class SynchExample1 {  
  
    static class Counter {  
        public static int count = 0;  
        public static void increment() {  
            count++;  
        }  
    }  
  
    static class MyRunnable implements Runnable {  
        @Override  
        public void run() {  
            for (int i = 0; i < 10000; i++)  
                Counter.increment();  
        }  
    }  
}
```

## 6.7 Monitors

```
public static void main(String[] args) throws Exception {  
    Thread[] threads = new Thread[5];  
    for (int i = 0; i < threads.length; i++) {  
        threads[i] = new Thread(new MyRunnable());  
        threads[i].start();  
    }  
    for (int i = 0; i < threads.length; i++)  
        threads[i].join();  
    System.out.println("counter = " + Counter.count);  
}  
  
}
```

## 6.7 Monitors

### ■ Java Synchronization Example 2:

```
public class SynchExample2 {  
  
    static class Counter {  
        public static int count = 0;  
        synchronized public static void increment() {  
            count++;  
        }  
    }  
  
    static class MyRunnable implements Runnable {  
        @Override  
        public void run() {  
            for (int i = 0; i < 10000; i++)  
                Counter.increment();  
        }  
    }  
}
```



## 6.7 Monitors

### ■ Java Synchronization Example 3:

```
public class SynchExample3 {  
  
    static class Counter {  
        private static Object object = new Object();  
        public static int count = 0;  
        public static void increment() {  
            synchronized (object) {  
                count++;  
            }  
        }  
    }  
}  
  
static class MyRunnable implements Runnable {  
    @Override  
    public void run() {  
        for (int i = 0; i < 10000; i++)  
            Counter.increment();  
    }  
}
```



## 6.7 Monitors

### ■ Java Synchronization Example 4:

```
public class SynchExample4 {  
  
    static class Counter {  
        public static int count = 0;  
        public void increment() {  
            synchronized (this) {  
                Counter.count++;  
            }  
        }  
    }  
}  
  
    static class MyRunnable implements Runnable {  
        Counter counter;  
        public MyRunnable(Counter counter) {  
            this.counter = counter;  
        }  
        @Override  
        public void run() {  
            for (int i = 0; i < 10000; i++)  
                counter.increment();  
        }  
    }  
}
```

## 6.7 Monitors

```
public static void main(String[] args) throws Exception {  
    Thread[] threads = new Thread[5];  
    for (int i = 0; i < threads.length; i++) {  
        threads[i] = new Thread(new MyRunnable(new Counter()));  
        threads[i].start();  
    }  
    for (int i = 0; i < threads.length; i++)  
        threads[i].join();  
    System.out.println("counter = " + Counter.count);  
}  
}
```

## 6.7 Monitors

### ■ Java Synchronization Example 5:

```
public static void main(String[] args) throws Exception {  
    Thread[] threads = new Thread[5];  
    Counter counter = new Counter();  
    for (int i = 0; i < threads.length; i++) {  
        threads[i] = new Thread(new MyRunnable(counter));  
        threads[i].start();  
    }  
    for (int i = 0; i < threads.length; i++)  
        threads[i].join();  
    System.out.println("counter = " + Counter.count);  
}
```



## 6.8 Liveness

### ■ Liveness

- Two criteria for the CSP: the progress and bounded-waiting.
  - Semaphores and monitors cannot solve these requirements.
- **Liveness** refers to
  - a set of properties that a system must satisfy
  - to ensure that processes make progress during their execution cycle.
- Two situations that can lead to liveness failures.
  - *deadlock* and *priority inversion*.





## 6.8 Liveness

### ■ *Deadlock*

- a situation where two or more processes are *waiting indefinitely*
  - for an event that *can be caused only by* one of the *waiting process*.

$P_0$	$P_1$
wait(S);	wait(Q);
wait(Q);	wait(S);
.	.
.	.
.	.
signal(S);	signal(Q);
signal(Q);	signal(S);



## 6.8 Liveness

### ■ *Priority Inversion*

- A situation where a higher-priority processes have to wait
  - for a lower-priority one to finish the resource.
- It can arise when a *higher*-priority process
  - needs to *read or modify kernel data*
  - that are currently being accessed by a *lower*-priority process.
- Typically, priority inversion is avoided
  - by implementing a ***priority-inheritance*** protocol.
- All processes accessing resources needed by a higher-priority process
  - inherit the higher priority
  - until they releases that resources.

*Any Questions?*

