



Chapter 8.

Deadlocks



**Operating System
Concepts (10th Ed.)**





8.1 System Model

- A **deadlock** is
 - a situation in which
 - every process in *a set of processes* is waiting for
 - an event that can be caused only *by another process in the set*.
 - a situation in which
 - a *waiting thread* (or process) can *never again change state*,
 - because the *resources* it has requested
 - are held *by other waiting threads* (or processes).



8.1 System Model

- Let us consider a system
 - consisting of *a finite number of resources*
 - to be distributed among a number of *competing threads*.
 - **Resource types** consist of
 - some number of **identical instances**.
 - e.g., CPU cycles, files, and I/O devices(such as printers, drives, etc.)
 - If a thread requests an *instance* of a *resource type*,
 - the allocation of **any instance** should *satisfy* the request.
 - A thread may utilize a resource as follows:
 - **Request** – **Use** – **Release**.



8.2 Deadlock in Multithreaded Applications

■ How can a deadlock occur?

```
/* thread_one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);

    pthread_exit(0);
}
```

```
pthread_mutex_t first_mutex;
pthread_mutex_t second_mutex;

pthread_mutex_init(&first_mutex, NULL);
pthread_mutex_init(&second_mutex, NULL);
```

```
/* thread_two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /**
     * Do some work
     */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);

    pthread_exit(0);
}
```

Figure 8.1 *Deadlock example.*



8.3 Deadlock Characterization

■ Four Necessary Conditions:

1. Mutual Exclusion:

- At least one resource is *held* in a *non-sharable* mode.

2. Hold and Wait:

- A thread *holds* at least one resource and *waiting* to acquire additional resources *held by* other threads.

3. No preemption:

- Resources *cannot* be *preempted*.

4. Circular Wait:

- A set of waiting threads exist such that the *dependency* graph of waiting is *circular*.



8.3 Deadlock Characterization

■ Resource-Allocation Graph:

- is a *directed graph* to describe deadlocks more precisely.
- consists of a set of vertices V and a set of edges E .
- Two different node types of V :
 - $T = \{T_1, T_2, \dots, T_n\}$: the set of all the *active threads* in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$: the set of all the *resource types* in the system.
- A directed edge: $T_i \rightarrow R_j$ (**request edge**)
 - signifies that a thread T_i has *requested* an instance of R_j .
- A directed edge: $R_j \rightarrow T_i$ (**assignment edge**)
 - signifies that an instance of R_j has been *allocated to* a thread T_i .



8.3 Deadlock Characterization





Figure 8.3 *Resource-allocation graph for program in Figure 8.1.*



8.3 Deadlock Characterization




- $T = \{T_1, T_2, T_3\}$
- $R = \{R_1, R_2, R_3, R_4\}$
- $E = \{T_1 \rightarrow R_1, T_2 \rightarrow R_3, R_1 \rightarrow T_2, R_2 \rightarrow T_2, R_2 \rightarrow T_1, R_3 \rightarrow T_3\}$

Figure 8.4 Resource-allocation graph.



8.3 Deadlock Characterization




- Two cycles exist in this graph.

$$\begin{aligned}
 &T_1 \rightarrow R_1 \rightarrow T_2 \rightarrow R_3 \rightarrow T_3 \rightarrow R_2 \rightarrow T_1 \\
 &T_2 \rightarrow R_3 \rightarrow T_3 \rightarrow R_2 \rightarrow T_2
 \end{aligned}$$

Figure 8.5 Resource-allocation graph with a deadlock.



8.3 Deadlock Characterization



- One cycle exists in this graph.

$$T_1 \rightarrow R_1 \rightarrow T_3 \rightarrow R_2 \rightarrow T_1$$

Figure 8.6 Resource-allocation graph with a cycle but no deadlock.



8.3 Deadlock Characterization

- An important observation:
 - If a *resource-allocation graph* does *not have a cycle*,
 - then the system is **not** in a deadlocked state.
 - If a *resource-allocation graph* *has a cycle*,
 - then the system **may** or **may not** be in a deadlocked state.

8.3 Methods for Handling Deadlocks

- Three ways of dealing with the Deadlock Problem:
 - *Ignore* the problem altogether
 - and pretend that deadlocks never occur in the system.
 - Use a protocol to *prevent* or *avoid* deadlocks,
 - ensuring that the system will *never enter* a deadlocked state.
 - Deadlock Prevention (§8.5)
 - *Deadlock Avoidance (§8.6): Banker's Algorithm*
 - Allow the system to enter a deadlocked state,
 - then *detect* it, and *recover* it.
 - Deadlock Detection (§8.7)
 - Recovery from Deadlock (§8.8)



8.5 Deadlock Prevention

■ Deadlock Prevention:

- For a deadlock to occur,
 - each of the *four necessary conditions* must hold.
 - Hence, we can **prevent** the occurrence of a deadlock,
 - by ensuring that *at least one* of these conditions cannot hold.
1. Mutual Exclusion
 2. Hold and Wait
 3. No Preemption
 4. Circular Wait



8.5 Deadlock Prevention

■ ***Mutual Exclusion***

- At least one resource must be non-sharable.
- In general, it *cannot be applied to* most applications.
 - some resources are *intrinsically* non-sharable.
 - e.g., a mutex lock cannot be shared by several threads.

■ ***Hold and Wait***

- We can guarantee that, whenever a thread requests a resource,
 - it does not hold any other resources.
- It is *impractical* for most applications.



8.6 Deadlock Avoidance

■ ***No preemption***

- We can use a protocol to ensure that there should be *preemption*.
- If a thread is holding some resources and requests another resources
 - that cannot be immediately allocated to it.
 - then, all resources the thread is currently holding are *preempted*.
- The preempted resources are added to the list of resources
 - for which the threads are waiting.
- The thread will be restarted
 - only when it can regain its *old resources* as well as *new ones*.
- *cannot generally be applied* to most applications.



8.6 Deadlock Avoidance

- ***Circular Wait***: sometimes *practical*.
 - Impose a *total ordering* of all resource types
 - and to *require* that each thread requests
 - resources in an *increasing order* of enumeration.
 - It is *provable* that these two protocols are used,
 - then the circular-wait condition *cannot hold*.
 - Note that, however,
 - imposing a lock ordering does not guarantee deadlock prevention,
 - if locks can be acquired dynamically.



8.6 Deadlock Avoidance

```
void transaction(Account from, Account to, double amount)
{
    mutex lock1, lock2;
    lock1 = get_lock(from);
    lock2 = get_lock(to);

    acquire(lock1);
    acquire(lock2);

    withdraw(from, amount);
    deposit(to, amount);

    release(lock2);
    release(lock1);
}

transaction(checking_account, savings_account, 25.0)

transaction(savings_account, checking_account, 50.0)
```

Figure 8.7 *Deadlock example with lock ordering.*



8.6 Deadlock Avoidance

- 8.17** In Section 8.5.4, we described a situation in which we prevent deadlock by ensuring that all locks are acquired in a certain order. However, we also point out that deadlock is possible in this situation if two threads simultaneously invoke the `transaction()` function. Fix the `transaction()` function to prevent deadlocks.



8.6 Deadlock Avoidance

- The **Demerits** of the Deadlock Prevention:
 - It prevents deadlocks by limiting how requests can made,
 - ensuring that *at least one of the necessary conditions* cannot occur.
 - However, possible side effects of preventing deadlocks are
 - *low device utilization* and *reduced system throughput*.



8.6 Deadlock Avoidance

■ Deadlock Avoidance:

- Let the system to decide for each request whether or not
 - the thread should *wait* in order to *avoid* a possible *future deadlock*.
- It requires *additional information* about
 - *how resources are to be requested*.
- For example, in a system with resources R_1 and R_2 ,
 - A thread P will request first R_1 and then R_2 before releasing them.
 - A thread Q will request R_2 then R_1 .



8.6 Deadlock Avoidance

- Given a *a priori* information,
 - it is possible to *construct an algorithm* that
 - ensures the system will *never enter* a deadlocked state.
 - Let the **maximum number** of resources of each type that it may need.
 - Let the **state** of resource allocation be
 - the number of **available** and **allocated** resources
 - and the **maximum demands** of the threads.



8.6 Deadlock Avoidance

■ Safe State:

- A state is *safe*
 - if the system *can allocate* resources to each thread (up to its *maximum*)
 - *in some order* and *still avoid* a deadlock.
- A system is in a safe state *if only if* there exists a **safe sequence**.
- A sequence of threads $\langle T_1, T_2, \dots, T_n \rangle$ is a *safe sequence*,
 - if, for each thread T_i , the resources that T_i can still request
 - can be satisfied by the currently available resources + resources held by all T_j , with $j < i$.



8.6 Deadlock Avoidance

- Basic facts:
 - A *safe state* is *not* a *deadlocked state*.
 - Conversely, a deadlocked state is an unsafe state.
 - However, not all unsafe states are deadlocks,
 - an unsafe state **may** lead to a deadlock.




Figure 8.8 *Safe, unsafe, and deadlocked state spaces.*



8.6 Deadlock Avoidance

- Given the concept of a **safe state**:
 - we can define an avoidance algorithm that
 - ensure that the system will *never enter* a deadlocked state.
 - The idea is simply to ensure that
 - the system will ***always remain in a safe state***.
 - *Initially*, the system is *in a safe state*.
 - Whenever a thread *requests* a *resource* that is currently *available*,
 - the system ***decides*** whether the resource can be ***allocated or not***.
 - The request is ***granted***
 - *if and only if* the allocation *leaves* the system *in a safe state*.



8.6 Deadlock Avoidance

- Revisit the Resource-Allocation Graph:
 - Suppose that a system has *only one instance* of each resource type.
 - Then, introduce a new type of edge, called a *claim edge*.
 - A *claim edge*: $T_i \rightarrow R_j$ indicates that
 - a thread *may request* a resource *at some time in the future*.
 - Then we can check for the safety
 - by a *cycle-detection* algorithm in a directed graph.
 - If *no cycle* exists, the request can be *granted* immediately,
 - since the resource allocation will *leave* the system in a *safe state*.
 - If a *cycle* is *detected*, then the request *cannot be granted*,
 - since the resource allocation will *put* the system in an *unsafe state*.



8.6 Deadlock Avoidance




Figure 8.9 *Resource-allocation graph for deadlock avoidance.*




Figure 8.10 *An unsafe state in a resource-allocation graph.*



8.6 Deadlock Avoidance

- Banker's Algorithm:
 - RAG is *not applicable* to a resource allocation system
 - with *multiple instances* of each resource type.
 - Banker's algorithm is *applicable* to such a system
 - but *less efficient* and *more complicated* than the RAG.
 - Why Banker's?
 - the *bank* never allocates its available cash in such a way that
 - it could no longer satisfy the needs of all its customers.



8.6 Deadlock Avoidance

■ Data structures:

- Let n be the number of *threads* in the system
 - and let m be the number of *resource types*.
- **Available**: A *vector* indicates the number of *available resource types*.
- **Max**: A *matrix* defines the *maximum demand* of each thread.
- **Allocation**: A *matrix* defines the number of resources of each type *currently allocated* to each thread.
- **Need**: A *matrix* indicates the *remaining resource need* of each thread.



8.6 Deadlock Avoidance

■ Data Structures:

- ***Available*** $[m]$:
 - if $Available[j] == k$, then k instances of R_j are available.
- ***Max*** $[n \times m]$:
 - if $Max[i][j] == k$, then T_i may request at most k instances of R_j .
- ***Allocation*** $[n \times m]$:
 - if $Allocation[i][j] == k$, then T_i is currently allocated k instances of R_j .
- ***Need*** $[n \times m]$:
 - if $Need[i][j] == k$, then T_i may need k more instances of R_j .



8.6 Deadlock Avoidance

■ Safety Algorithm:

1. Let *Work* and *Finish* be vectors of length m and n , respectively. Initialize *Work* = *Available* and *Finish*[i] = *false* for $i = 0, 1, \dots, n - 1$.
2. Find an index i such that both
 - a. *Finish*[i] == *false*
 - b. $Need_i \leq Work$If no such i exists, go to step 4.
3. *Work* = *Work* + *Allocation* _{i}
Finish[i] = *true*
Go to step 2.
4. If *Finish*[i] == *true* for all i , then the system is in a safe state.



8.6 Deadlock Avoidance

■ Resource-Request Algorithm:

1. If $Request_i \leq Need_i$, go to step 2. Otherwise, raise an error condition, since the thread has exceeded its maximum claim.
2. If $Request_i \leq Available$, go to step 3. Otherwise, T_i must wait, since the resources are not available.
3. Have the system pretend to have allocated the requested resources to thread T_i by modifying the state as follows:

$$\begin{aligned} Available &= Available - Request_i \\ Allocation_i &= Allocation_i + Request_i \\ Need_i &= Need_i - Request_i \end{aligned}$$



8.6 Deadlock Avoidance

■ An illustrative example:

- a set of five threads: $T = \{T_0, T_1, T_2, T_3, T_4\}$
- a set of three resource types: $R = \{A, B, C\}$
- the number of instances of each resource types: $A = 10, B = 5, C = 7$
- the snapshot representing the current state of the system:

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	A B C	A B C	A B C
T_0	0 1 0	7 5 3	3 3 2
T_1	2 0 0	3 2 2	
T_2	3 0 2	9 0 2	
T_3	2 1 1	2 2 2	
T_4	0 0 2	4 3 3	



8.6 Deadlock Avoidance

- Note that $Need[i][j] = Max[i][j] - Allocation[i][j]$.

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>		<u>Need</u>
	A B C	A B C	A B C		A B C
T_0	0 1 0	7 5 3	3 3 2	T_0	7 4 3
T_1	2 0 0	3 2 2		T_1	1 2 2
T_2	3 0 2	9 0 2		T_2	6 0 0
T_3	2 1 1	2 2 2		T_3	0 1 1
T_4	0 0 2	4 3 3		T_4	4 3 1



8.6 Deadlock Avoidance

- Now we claim that the system is currently in a safe state.
 - In deed, the sequence $\langle T_1, T_3, T_4, T_2, T_0 \rangle$ satisfies the safety criteria.

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
T_0	0 1 0	7 4 3	3 3 2
T_1	2 0 0	1 2 2	
T_2	3 0 2	6 0 0	
T_3	2 1 1	0 1 1	
T_4	0 0 2	4 3 1	



8.6 Deadlock Avoidance

- When a **new request** is submitted:
 - Suppose that T_1 requests one instance of A and two instances of C.
 - $Request_1 = (1, 0, 2)$,
 - Decide whether this request should be *granted or not*.

$$Request_1 \leq Available$$

$$(1, 0, 2) \leq (3, 3, 2)$$

$$(3, 3, 2) - (1, 0, 2) = (2, 3, 0)$$

$$(1, 2, 2) - (1, 0, 2) = (0, 2, 0)$$

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
T_0	0 1 0	7 4 3	2 3 0
T_1	3 0 2	0 2 0	
T_2	3 0 2	6 0 0	
T_3	2 1 1	0 1 1	
T_4	0 0 2	4 3 1	



8.6 Deadlock Avoidance

- Now, determine whether this new system state is safe.
 - Safety algorithm finds that $\langle T_1, T_3, T_4, T_0, T_2 \rangle$ satisfies the safety.

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
T_0	0 1 0	7 4 3	2 3 0
T_1	3 0 2	0 2 0	
T_2	3 0 2	6 0 0	
T_3	2 1 1	0 1 1	
T_4	0 0 2	4 3 1	



8.6 Deadlock Avoidance

- Now, determine with a request of $(3, 3, 0)$ by T_4 .
 - $Request_4 = (3, 3, 0)$,

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
T_0	0 1 0	7 4 3	2 3 0
T_1	3 0 2	0 2 0	
T_2	3 0 2	6 0 0	
T_3	2 1 1	0 1 1	
T_4	0 0 2	4 3 1	



8.6 Deadlock Avoidance

- How about a request of $(0, 2, 0)$ by T_0 ?
 - $Request_0 = (0, 2, 0)$,

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
T_0	0 1 0	7 4 3	2 3 0
T_1	3 0 2	0 2 0	
T_2	3 0 2	6 0 0	
T_3	2 1 1	0 1 1	
T_4	0 0 2	4 3 1	



8.7 Deadlock Detection

- Deadlock Detection:
 - If a system does not *prevent* or *avoid* the deadlock,
 - then a deadlock situation may occur.
 - In this environment, the system may provide:
 - An algorithm that examines the state of the system to *determine whether* a deadlock has *occurred*.
 - An algorithm to *recover* from the deadlock.



8.6 Deadlock Avoidance

- **Single Instance** of Each Resource Type
 - Maintain a *wait-for graph*,
 - a variant of the resource-allocation graph.
 - *Periodically, invoke an algorithm* that
 - searches for a cycle in the *wait-for* graph.



8.6 Deadlock Avoidance

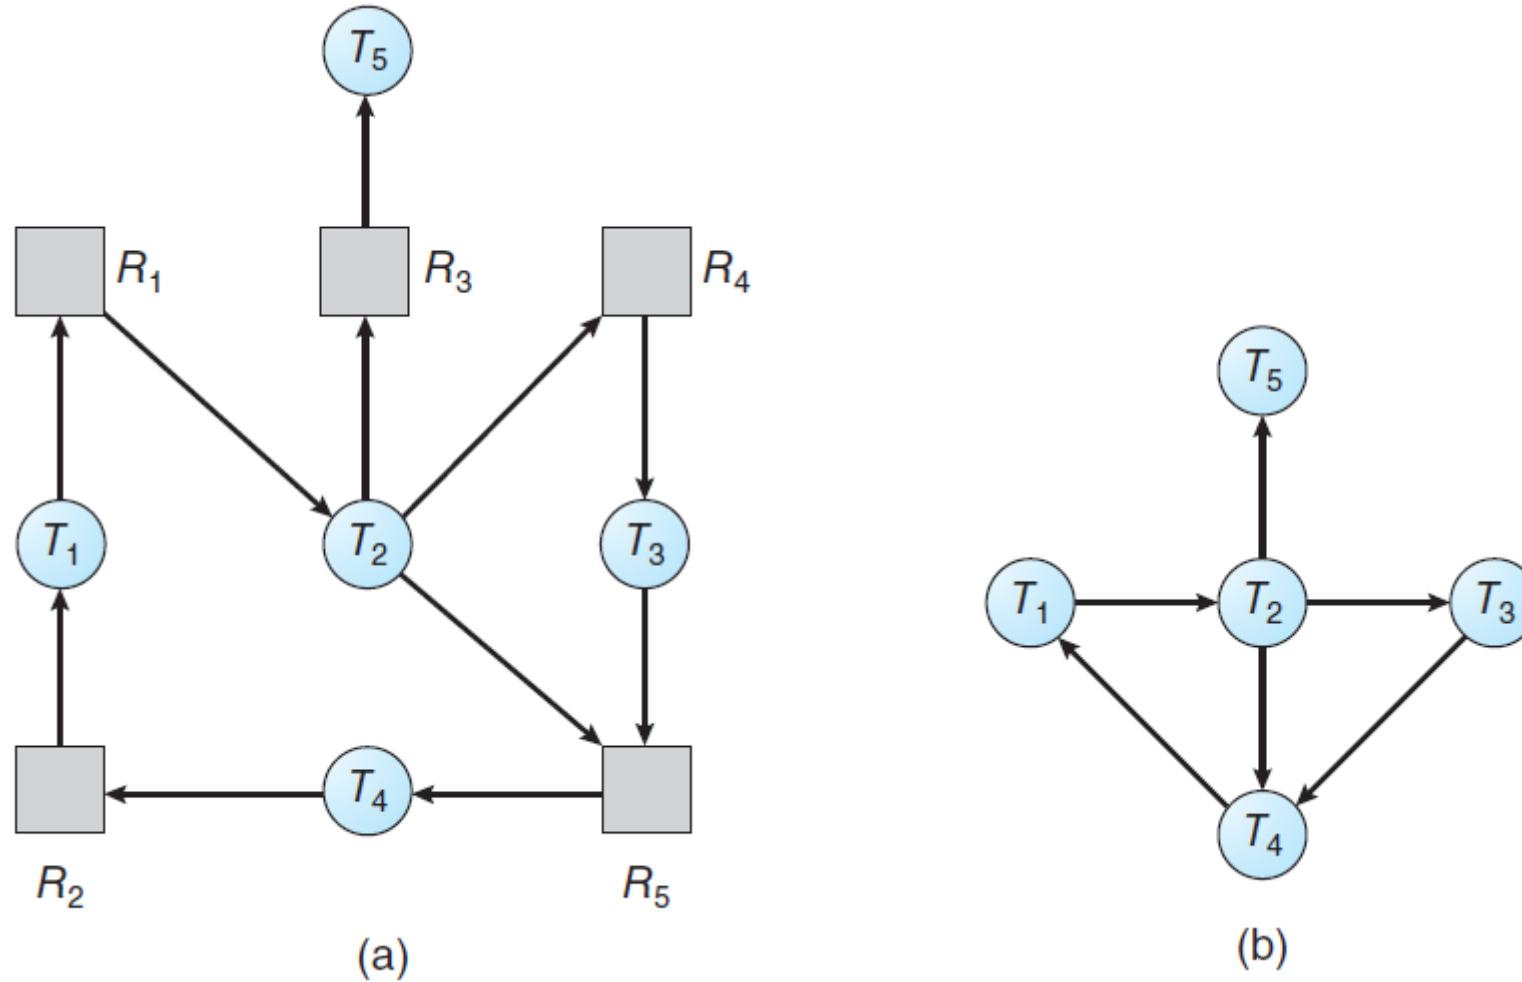


Figure 8.11 (a) Resource-allocation graph. (b) Corresponding wait-for graph.



8.7 Deadlock Detection

- **Several Instances** of a Resource Type:
 - The wait-for graph is not applicable to a system
 - with *multiple instances* of each resource type.
 - We can design a deadlock detection algorithm
 - that is *similar to* those used in *the banker's algorithm*.



8.7 Deadlock Detection

- Data Structures:
 - *Available*[m]:
 - *Allocation*[$n \times m$]:
 - *Request*[$n \times m$]: indicates the *current request* of each thread.
 - if *Request*[i][j] == k , then T_i is requesting k more instances of R_j .



8.7 Deadlock Detection

■ Detection Algorithm:

1. Let *Work* and *Finish* be vectors of length m and n , respectively. Initialize *Work* = *Available*. For $i = 0, 1, \dots, n-1$, if $Allocation_i \neq 0$, then $Finish[i] = false$. Otherwise, $Finish[i] = true$.
2. Find an index i such that both
 - a. $Finish[i] == false$
 - b. $Request_i \leq Work$If no such i exists, go to step 4.
3. $Work = Work + Allocation_i$
 $Finish[i] = true$
Go to step 2.
4. If $Finish[i] == false$ for some $i, 0 \leq i < n$, then the system is in a deadlocked state. Moreover, if $Finish[i] == false$, then thread T_i is deadlocked.



8.7 Deadlock Detection

- An illustrative example:
 - a set of five threads: $T = \{T_0, T_1, T_2, T_3, T_4\}$
 - a set of three resource types: $R = \{A, B, C\}$
 - the number of instances of each resource types: $A = 7, B = 2, C = 6$.
 - the snapshot representing the current state of the system:

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
T_0	0 1 0	0 0 0	0 0 0
T_1	2 0 0	2 0 2	
T_2	3 0 3	0 0 0	
T_3	2 1 1	1 0 0	
T_4	0 0 2	0 0 2	



8.7 Deadlock Detection

- Now we claim that the system is *not in a deadlocked state*.
 - the sequence $\langle T_0, T_2, T_3, T_1, T_4 \rangle$ results in $Finish[i] == true$ for all i .

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
T_0	0 1 0	0 0 0	0 0 0
T_1	2 0 0	2 0 2	
T_2	3 0 3	0 0 0	
T_3	2 1 1	1 0 0	
T_4	0 0 2	0 0 2	



8.7 Deadlock Detection

- Now we claim that the system is *now deadlocked*.
 - a deadlock exists, consisting of threads T_1, T_2, T_3 , and T_4 .

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
T_0	0 1 0	0 0 0	0 0 0
T_1	2 0 0	2 0 2	
T_2	3 0 3	0 0 1	
T_3	2 1 1	1 0 0	
T_4	0 0 2	0 0 2	



8.8 Recovery from Deadlock

- *When* should we invoke the detection algorithm?
 - How *often* is a deadlock likely to occur?
 - *more frequent* deadlocks, *more frequent* deadlock detections.
 - How *many* threads will be affected by deadlock when it happens?
 - the number of threads involved in the deadlock cycle *may grow*.
- Invoking *for every request* .vs. invoking *at defined intervals*.
 - Note that there is a considerable overhead in computation time.
 - However, there may be many cycles in the resource graph,
 - if the detection algorithm is invoked at arbitrary points in time.



8.4 Recovery from Deadlock

- When a detection algorithm determines a deadlock exists,
 - *inform the operator* that a deadlock has occurred.
 - or let the system *recover from* the deadlock *automatically*.
 - Process and Thread Termination
 - Resource Preemption



8.4 Recovery from Deadlock

■ Deadlock Recovery

- Process and Thread Termination:
 - Abort *all* deadlocked processes.
 - Abort *one* process *at a time* until the deadlock cycle is eliminated.
- Resource Preemption:
 - Selecting a victim: consider the order of preemption to *minimize cost*.
 - Rollback: *roll back* the process to some *safe state* and *restart* it.
 - Starvation: *picked as a victim* only a finite number of times.

Any Questions?

