



Chapter 10.

Virtual Memory



**Operating System
Concepts (10th Ed.)**





10.1 Background

■ *Virtual Memory*

- a technique to allow the execution of processes
 - that are not completely in memory
 - so, programs can be larger than physical memory.
- abstracts main memory into an extremely large array of storage,
 - separating logical memory from physical memory.
- provides an efficient mechanism
 - for sharing files and libraries and process creation.



10.1 Background




Figure 10.1 *Diagram showing virtual memory that is larger than physical memory.*



10.1 Background

■ *Virtual Address Space*

- the logical (or virtual) view of how a process is stored in memory.
- Typically, begins a certain logical address, to say, address 0,
- and exists in contiguous memory.




Figure 10.2 *Virtual address space of a process in memory.*



10.1 Background

5

■ Virtual Memory

- allows *files* and *memory* to be shared by two or more processes
 - through *page sharing*.




Figure 10.3 *Shared library using virtual memory.*



10.2 Demand Paging

- Consider how an executable program
 - might be loaded from secondary storage into memory.
 - One option is to load the entire program in physical memory
 - The *demand paging* is an alternative strategy
 - to load pages only as they are needed.
 - commonly used in virtual memory systems.
 - With *demand-paged virtual memory*,
 - pages are loaded only when they are *demanded* during execution.



10.2 Demand Paging

- Basic Concepts of **the Demand Paging**:
 - While a process is executing,
 - some pages will be *in memory* and some will be *in secondary storage*.
 - To *distinguish* between these two situations,
 - the *valid-invalid bit* scheme can be used.
 - *valid*: the page is both legal and in memory.
 - *invalid*: the either is not valid or currently in secondary storage.



10.2 Demand Paging




Figure 10.4 *Page table when some pages are not in main memory.*



10.2 Demand Paging

- The procedure for handling the *Page Fault*:
 1. *Check an internal table* for the process to determine
 - whether the reference was *valid* or *invalid* memory access.
 2. If the reference was valid, terminate the process,
 - or valid but page fault, we now *page it in*.
 3. Find a *free frame* (by taking one from the free-frame list)
 4. Schedule a secondary storage operation
 - to *read the desired page* into the newly allocated frame.
 5. When the storage read is complete, *modify the internal table and*
 - *the page table* to indicate that the page is now in memory.
 6. *Restart the instruction* that was interrupted by the trap.



10.2 Demand Paging




Figure 10.5 Steps in handling a page fault.



10.2 Demand Paging

- **Pure** Demand Paging:
 - *never bring a page* into memory *until it is required*.
 - With the scheme of pure demand paging,
 - we can start executing a process with *no pages* in memory.
 - When the OS sets the instruction pointer
 - to the *first instruction* of the process with a *page fault*,
 - the page of the process would be paged in.



10.2 Demand Paging

■ *Locality of Reference:*

- If a program accesses *several new pages* with *each instruction*,
 - to say, *one page* for *the instruction* and *many pages* for *data*,
 - possibly causes *multiple page faults per instruction*.
- Fortunately, analysis running processes
 - show that this behavior is exceedingly unlikely.
- Programs tend to have *the locality of reference*,
 - which results in *reasonable performance* from *demand paging*.



10.2 Demand Paging

■ An Example of Program Structure (§10.9.5)

- Let the *page size* to be 128 and an *array* to be 128×128 ,
 - then compare the following two codes.

```
int i, j;  
int[128][128] data;  
  
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i][j] = 0;
```

```
int i, j;  
int[128][128] data;  
  
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        data[i][j] = 0;
```

- Careful selection of data structures and programming structures
 - can increase *the locality of code or data*
 - hence, *lower the page-fault rate and enhance the system performance.*



10.2 Demand Paging

- Hardware Support to Demand Paging:
 - *Page table*:
 - has the ability to mark valid or invalid. (with a valid-invalid bit)
 - *Secondary memory*: (**=swap space**)
 - holds those pages that are not present in main memory.
 - usually a high-speed disk or NVM device



10.2 Demand Paging

■ Instruction Restart:

- A crucial requirement for demand paging.
 - the ability to restart any instruction after page fault.
- When the page fault occurs,
 - the state of interrupted process (registers, condition code, instruction counter, etc.) is saved.
- Therefore, *restart* the process in *exactly the same place and state*.
- If a *page fault* occurs *on the instruction fetch*,
 - restart by fetching the instruction again.
- If a *page fault* occurs while we are *fetching an operand*,
 - fetch and decode the instruction again and then fetch the operand.



10.2 Demand Paging

- As a worst-case example:
 - `ADD A, B, C`; three address instruction, adding A and B into C.
 1. Fetch and decode the instruction (ADD)
 2. Fetch A
 3. Fetch B
 4. ADD A and B
 5. Store the sum in C



10.2 Demand Paging

■ Free Frame List:

- When a page fault occurs,
 - OS must bring the desired page from secondary storage into memory.
- To resolve page faults, OS maintains the *free frame list*:
 - a pool of free frames for satisfying such requests.
- Free frames must also be allocated
 - when the stack or heap segments from a process expand.



Figure 10.6 *List of free frames.*



10.2 Demand Paging

- Performance of Demand Paging:
 - How to compute the *effective access time* for a *demand-paged* memory?
 - Let *ma* to denote the *memory-access time*.
 - Let *p* be the *probability* of a *page fault*.
 - $EAT = (1 - p) \times ma + p * (\text{page fault time})$.
 - How much time is needed to service a page fault?
 - three major activities:
 - Service the page-fault interrupt.
 - Read in the page.
 - Restart the process.



10.2 Demand Paging

- Consider a system:
 - with an average page-fault service time of *8 milliseconds*
 - and a memory access time of *200 nanoseconds*,
 - $EAT = (1 - p) \times 200 + p \times 8,000,000$
 $= 200 + 7,999,800 \times p$
 - If one access out of 1,000 causes a page fault ($p = 0.001$),
 - $EAT = 200 + 7999.8 = 8199.8$ *nanoseconds*
 $\cong 8.2$ *microseconds*



10.3 Copy-on-Write

■ Copy-on-Write:

- *copy* a shared page *only when* a process *writes* to a shared page.
- Recall the process creation with **fork()** and **exec()**.




Figure 10.7 *Before* process 1 modifies page C.

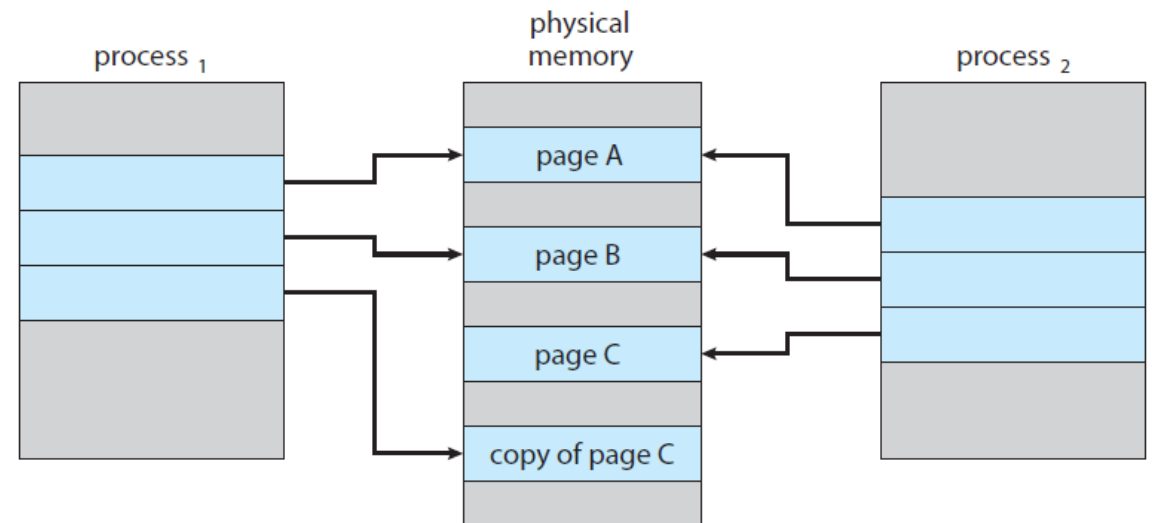


Figure 10.8 *After* process 1 modifies page C.



10.4 Page Replacement

- What happens if there is *no free frames*?
 - If we increase our degree of multiprogramming,
 - we are over-allocating memory.
 - If we have **40 frames** and run **6 processes**,
 - each of which is **10 pages** in size, but actually uses only **5 pages**.
 - Then, we manage to demand-paged system with 10 frames spared.
 - However, what if the processes *suddenly want to use all 10 pages*,
 - or need a huge buffer *consuming more pages than available* ones?



10.4 Page Replacement




Figure 10.9 *Need for page replacement.*



10.4 Page Replacement

■ Page Replacement:

- If no frame is free,
 - find one that is not currently being used and free it.
- Free a frame by writing its contents to swap space
 - and changing the page table
 - to indicate that the page is no longer in memory (invalid or dirty).
- Now, use the freed frame to hold page
 - for which the process faulted.



10.4 Page Replacement

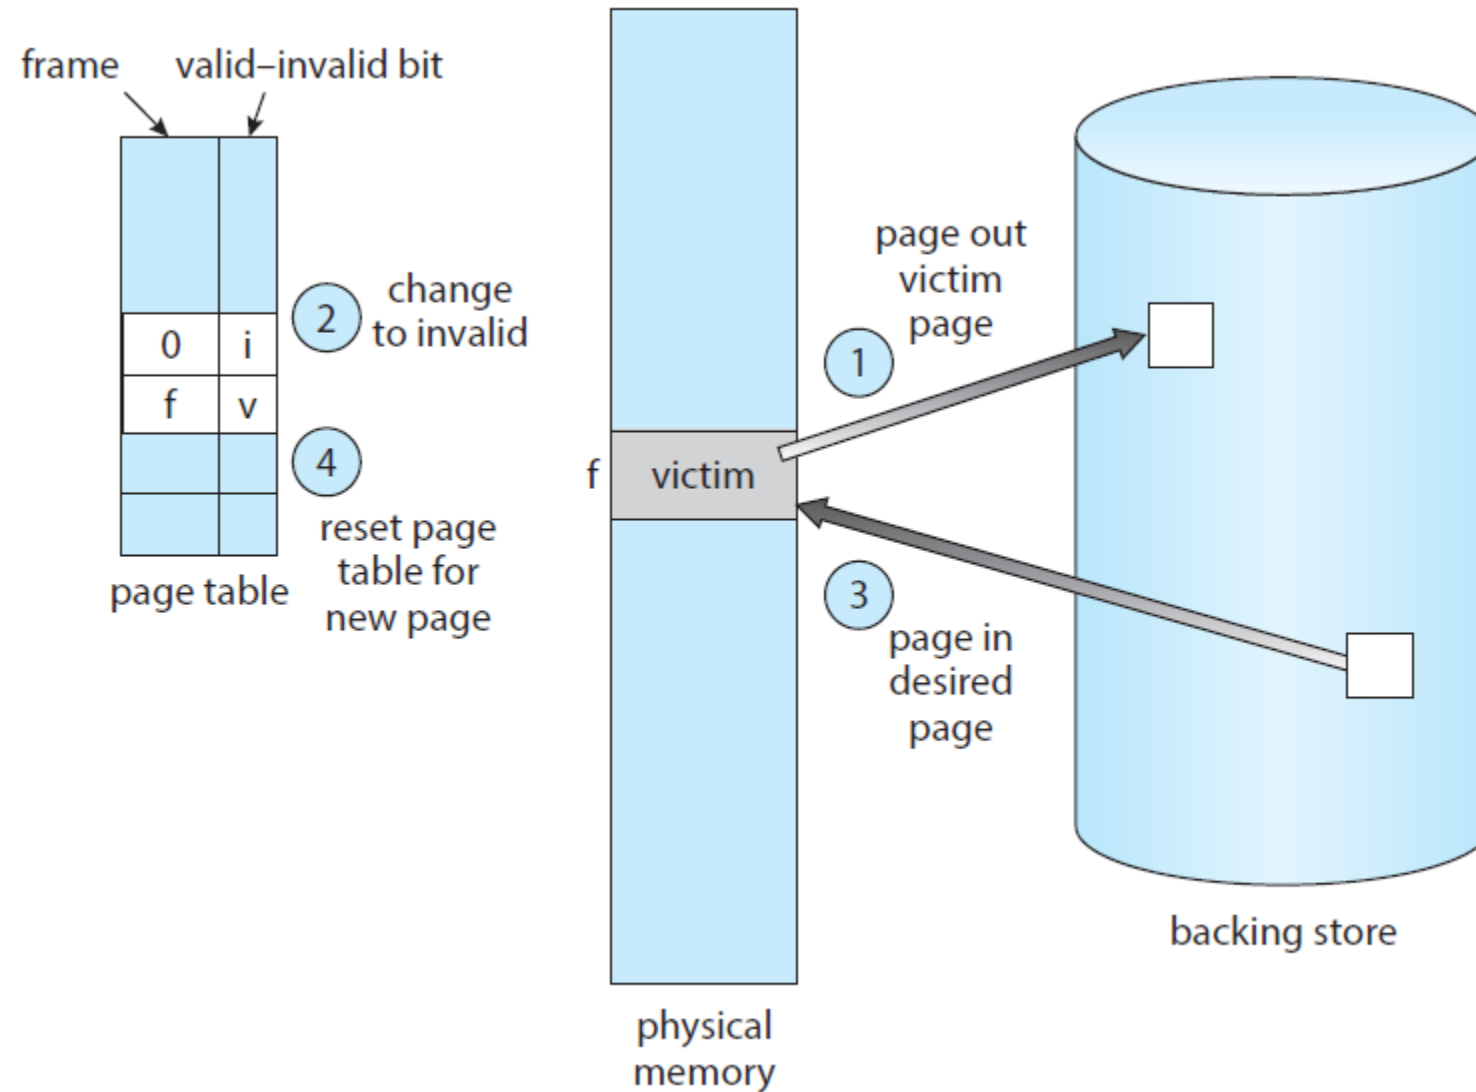


Figure 10.10 *Page replacement.*



10.4 Page Replacement

- Page Fault Service Routine incl. Page Replacement:
 1. Find the location of the desired page on secondary storage.
 2. Find *a free frame*:
 - If there is a free frame, use it.
 - If there is *no free frame*, use a *page-replacement algorithm* to select *a victim frame*.
 - Write the victim frame to secondary storage; change the page and frame tables accordingly.
 - Read the desired page into the *newly freed frame*
 - change the page and frame tables.
 - Continue the process from where the *page fault occurred*.



10.4 Page Replacement

- Two major problems to implement demand paging:
 - *Frame-allocation algorithm*:
 - how many frames to allocated to each process?
 - *Page-replacement algorithm*:
 - select the frames that are to be replaced.
 - Since the secondary storage I/O is *so expensive*,
 - even *slight improvements* in demand-paging methods
 - can yield *large gains* in system performance.



10.4 Page Replacement

- Evaluation of Page Replacement Algorithms:
 - **reference string**: a string of memory references.
 - Evaluate an algorithm by running it on a reference string
 - and computes *the number of page faults*. (**minimize it!**)
 - What about the number of page frames?
 - Obviously, the more frames, the less page faults.





Figure 10.11 Graph of page faults versus number of frames.



10.4 Page Replacement

- An example to evaluate algorithms:
 - reference string: **7 0 1 2 0 3 0 4 2 3 0 3 0 3 2 1 2 0 1 7 0 1**
 - in a memory with *three* frames.





10.4 Page Replacement

■ **FIFO** Page Replacement:

- FIFO: First-In-First-Out: the simplest algorithm.
- Choose *the oldest page* when a page must be replaced.
- There are *15 page faults* with our example.

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																
	0	0	0		2	2	4	4	4	0			0	0			7	7	7
		1	1		3	3	3	2	2	2			1	1			1	0	0
					1	0	0	0	3	3			3	2			2	2	1

page frames


Figure 10.12 *FIFO page-replacement algorithm.*



10.4 Page Replacement

■ *Belady's Anomaly:*

- The page-fault rate *may increase*
 - as the number of allocated frames *increases*.



- **reference string:**
1 2 3 4 1 2 5 1 2 3 4 5

Figure 10.13 Page-fault curve of FIFO replacement on a reference string.



10.4 Page Replacement

■ **Optimal** Page Replacement:

- Seeking for an *optimal* algorithm that has the *lowest page-fault rate*,
 - and *never suffers from* Belady's anomaly.
- OPT or MIN:
 - replace the page that *will not be used* for the *longest period of time*.
- OPT will guarantee the lowest possible page-fault rate.



10.4 Page Replacement

- The difficulty of implementing OPT:
 - There are *9 page faults* with our example.
 - OPT requires *future knowledge* of the reference string.
 - used mainly for comparison studies.




Figure 10.14 *Optimal page-replacement algorithm.*



10.4 Page Replacement

- Recall the Shortest-Job-First CPU scheduler.
 - The key distinction between the FIFO and the OPT:
 - looking *backward*: when a page was *brought in*?
 - looking *forward*: when a page to *be used*?
 - If we use the *recent past*
 - as an *approximation* of the *near future*,
 - then we can replace the page
 - that has *not been used* for the *longest period of time*.



10.4 Page Replacement

■ LRU Page Replacement:

- LRU: Least Recently Used
- Associates with each page *the time of that page's last use*, and
 - choose the page that has *not been used* for the *longest period* of time.
 - There are *12 page faults* with our example.

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0		1		1		1			
	0	0	0		0		0	0	3	3		3		0		0			
		1	1		3		3	2	2	2		2		2		7			

page frames

Figure 10.15 LRU page-replacement algorithm.



10.4 Page Replacement

- LRU policy
 - is considered to be *good* and is *often used*.
 - However, the problem to solve for the implementation of LRU is
 - to determine an order for the frames defined by the time of last use?
 - It may require substantial *hardware assistance*.
 - Two implementations are possible: *counter* and *stack*
 - LRU *does not suffer* from Belady's anomaly like OPT.



10.4 Page Replacement

- Two implementation methods for the LRU:
 - *Counter* implementation:
 - Whenever a page is referenced, copy the counter (or the clock).
 - Replace the page with the smallest value.
 - *Stack* implementation:
 - Keep a stack of page numbers.
 - Note that entries must be removed from the middle of the stack.




Figure 10.16 *Use of a stack to record the most recent page reference.*



10.4 Page Replacement

■ LRU-Approximation

- LRU needs hardware support,
 - however, many systems provide some help with a *reference bit*.
- *reference bit*:
 - initially 0, a bit associated with each page.
 - when a page is referenced, set to 1.
 - replace any with reference bit = 0 (if any).



10.4 Page Replacement

■ Second-Change Algorithm

- Use a FIFO replacement algorithm.
- However, inspect its reference bit, when a page has been selected.
 - If the value is 0, proceed to replace it.
 - If the value is 1, give the page a second change and move on to select the next FIFO page.
- When a page gets a second change,
 - its reference bit is cleared,
 - and its arrival time is reset to the current time.



10.4 Page Replacement

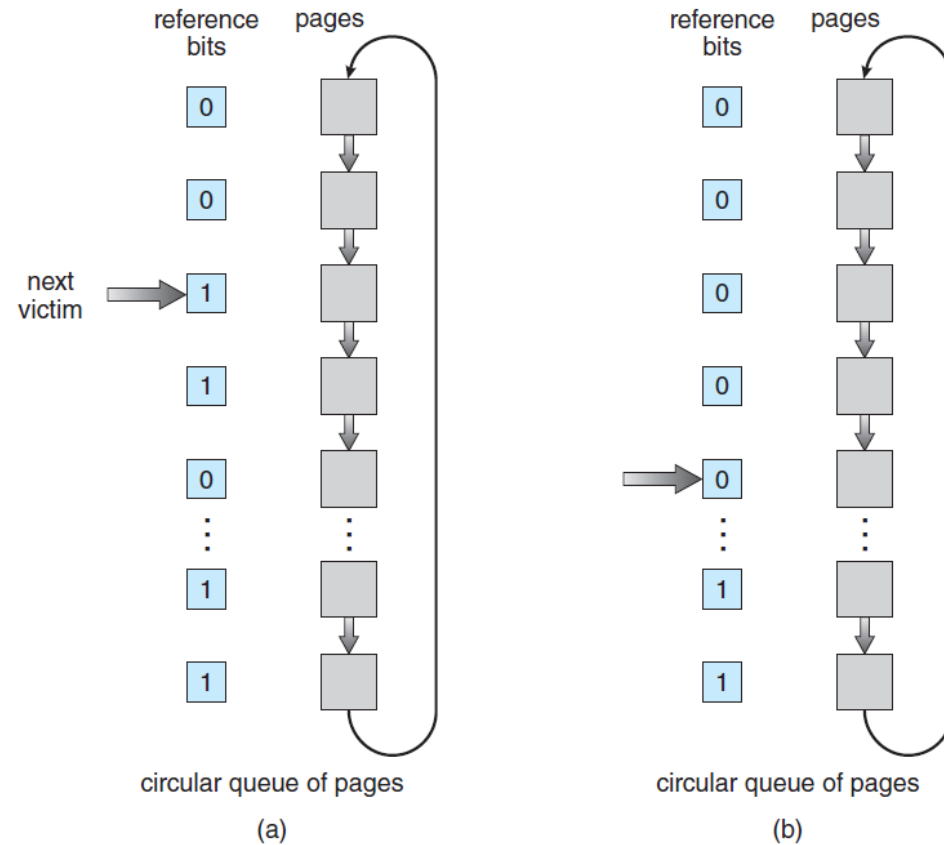


Figure 10.17 *Second-chance (clock) page-replacement algorithm.*



10.5 Allocation of Frames

- The Issues for Frame Allocation:
 - Consider a simple case of a system with 128 frames.
 - OS may take 35, leaving *93 frames* for the user process.
 - Using the *pure-demand-paging*,
 - 93 frames would be put on the *free-frame list*.
 - The first 93 page faults would get free frames.
 - 94th *page faults* would cause a *page replacement*.
 - Then, if we have *two processes*,
 - how do we *allocate 93 frames* to these *two processes*?



10.5 Allocation of Frames

- The strategies for frame allocation
 - Equal .vs. Proportional:
 - *equal allocation*: give every process an equal share.
 - *proportional allocation*: allocate according to the size of process.
 - Global .vs. Local:
 - *local replacement*: selects from *only its own* set of allocated frames.
 - *global replacement*: selects a replacement frame from the set of all the frames in the system.



10.6 Thrashing

■ Thrashing:

- a situation that a process is busy swapping pages in and out.
- If a process does *not* have *enough pages*,
 - the page-fault rate is very high.




Figure 10.20 Thrashing.



10.6 Thrashing

■ Working-Set Model:

- based on the assumption of *locality*.
- define the *working-set window* with a parameter Δ .
- The idea is to examine the most recent Δ page references.
 - **working-set**: the set of pages in the most recent Δ page references
- If a page is in active use,
 - it will be in the working set.
- If it is no longer being used,
 - it will drop from the working set time units after its last reference.



10.6 Thrashing




Figure 10.21 *Locality in a memory-reference pattern.*

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...




Figure 10.22 *Working-set model.*

Any Questions?

