


Chapter 11-15.

Storage Management



**Operating System
Concepts (10th Ed.)**





■ Mass-Storage

- non-volatile, secondary storage system of a computer
- usually, HDD(Hard Disk Drive) or NVM(Non-Volatile Memory).
- sometimes, magnetic tapes, optical disks, or cloud storage.
 - using the structure of RAID systems.



■ Hard Disk Drives:




Figure 11.1 *HDD moving-head disk mechanism.*



Ch.11 Mass-Storage Structure

- HDD Scheduling:
 - to *minimize* the access time (or *seek time*)
 - to *maximize* data transfer *bandwidth*.
- *seek time*:
 - the time for the device arm to move the heads to the cylinder containing the desired sector,
 - and the rotational latency is the additional time for the platter to rotate the desired sector to the head
- *disk bandwidth*:
 - the total number of bytes transferred, divided by the total time
 - between the first request and the completion of the last transfer.



■ FIFO Scheduling:

- intrinsically fair, but generally does not provide the fastest services.
- total *head movement* of **640 cylinders**.




Figure 11.6 FCFS disk scheduling.



■ SCAN Scheduling:

- the disk arm *starts* at *one end* of the disk and moves *toward the other end*,
 - servicing requests as it reaches each cylinder,
 - until it gets to the other end of the disk.
- At the other end,
 - the *direction* of head movement is *reversed*, and continue.



■ SCAN Scheduling:

- if the direction of head movement is moving towards 0.
- total *head movement* of **236 cylinders**.




Figure 11.7 SCAN disk scheduling.



- **C-SCAN** (Circular-SCAN) Scheduling:
 - a variant of SCAN designed to provide a *more uniform wait time*.
 - moves the head from one end of the disk to the other,
 - servicing requests along the way.
 - however, when the head reaches the other end,
 - *returns immediately to the beginning* of the disk
 - *without servicing* any requests on the return trip.
 - treats the cylinders as a *circular list*
 - that wraps around from the final cylinder to the first one.



■ C-SCAN Scheduling:

- total *head movement* of **183 cylinders**. (ignore from 199 to 0)




Figure 11.8 C-SCAN disk scheduling.



Ch.11 Mass-Storage Structure

■ Boot Block

- For a computer to start running, when it is powered up,
 - it must have an initial program to run.
- A bootstrap loader is stored in NVM flash memory,
 - and mapped to a known memory location.

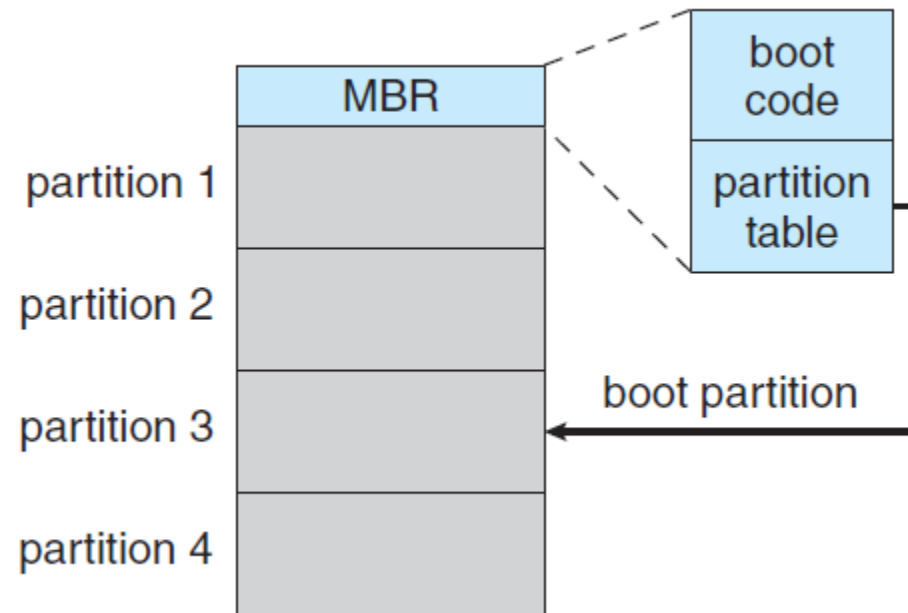


Figure 11.10 *Booting from a storage device in Windows.*



Ch.11 Mass-Storage Structure

- **RAID**: Redundant Arrays of Independent Disks
 - a collection of a variety of disk-organization techniques.
 - to improve the *rate* at which data can be read or written.
 - if the drives are operated *in parallel*.
 - to improve the *reliability* of data storage,
 - because *redundant information* can be stored on multiple devices.
 - thus, *failure* of one drive does *not* lead to *loss of data*.



Ch.11 Mass-Storage Structure

- *Redundancy*: Improvement of **Reliability**
 - The *chance of fail* in a set of *N disks*
 - is much *greater* than the chance in a *single* disk.
 - Suppose that *mean time between failures* (MTBF) = 100,000 hours.
 - MTBF in an array of 100 disks = $100,000 / 100 = 1,000$ hours.
 - The solution the problem of reliability: ***redundancy***.
 - the simplest: *mirroring* (duplicate all the drives)



Ch.11 Mass-Storage Structure

- *Parallelism*: Improvement in **Performance**
 - With multiple drives, we can improve the transfer rate
 - by *striping* data across the drives.
 - ***bit-level striping***: splitting the *bits* of each *byte*.
 - if we have 8 drives, we can write i bit of each byte to drive i .
 - ***block-level striping***: generalization to a number of drives.



Ch.11 Mass-Storage Structure

- RAID Levels:
 - *mirroring*: highly reliable, however, too *expensive*.
 - *striping*: highly efficient, however, *not* related to *reliability*.
 - How about *parity bit*?
 - set to **1**: the number of bits in the byte is *even*.
 - set to **0**: the number of bits in the byte is *odd*.
 - It enable us to detect an error if f one of the bits is damaged.
 - *RAID levels*: classify these schemes
 - according to *different cost-performance trade-offs*



Ch.11 Mass-Storage Structure

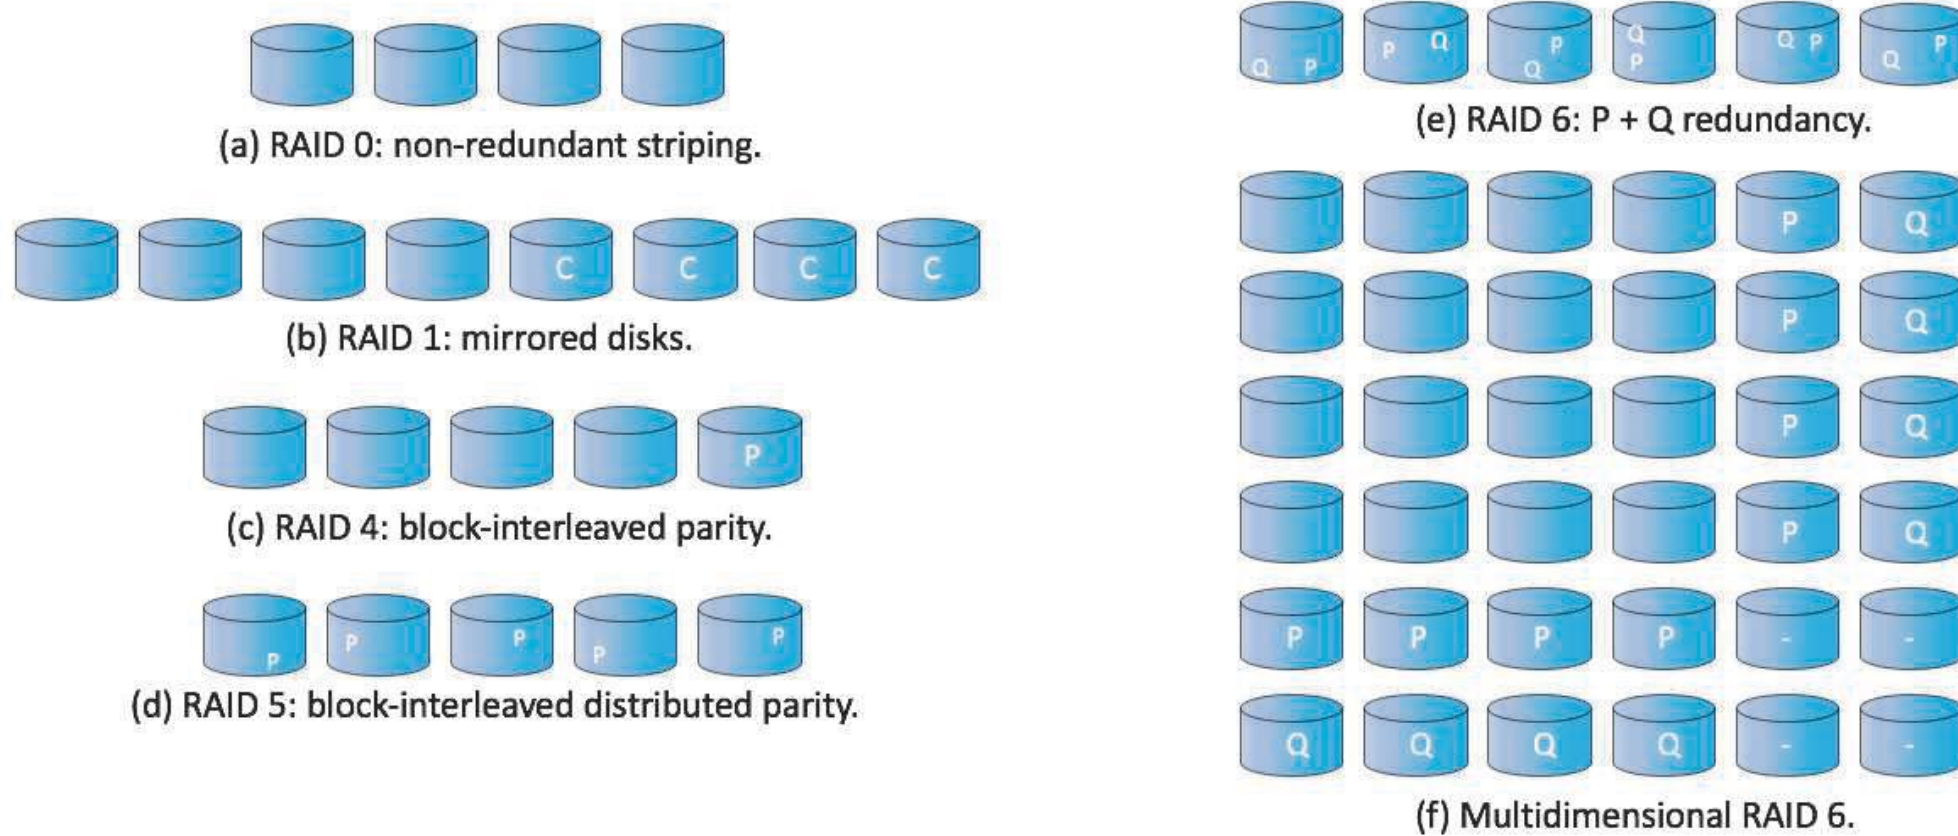




Figure 11.15 RAID levels.



Ch.11 Mass-Storage Structure



a) RAID 0 + 1 with a single disk failure.



b) RAID 1 + 0 with a single disk failure.

Figure 11.16 RAID 0 + 1 and 1 + 0 with a single disk failure.



Ch. 12 I/O Systems

- Two main jobs of a computer: I/O and computing.
 - In many cases, the main job is I/O,
 - for instance, web browsing, file editing, youtube, game, and so force.
 - The role of Operating System in I/O is
 - to manage and control I/O operations and I/O devices.



Ch. 12 I/O Systems

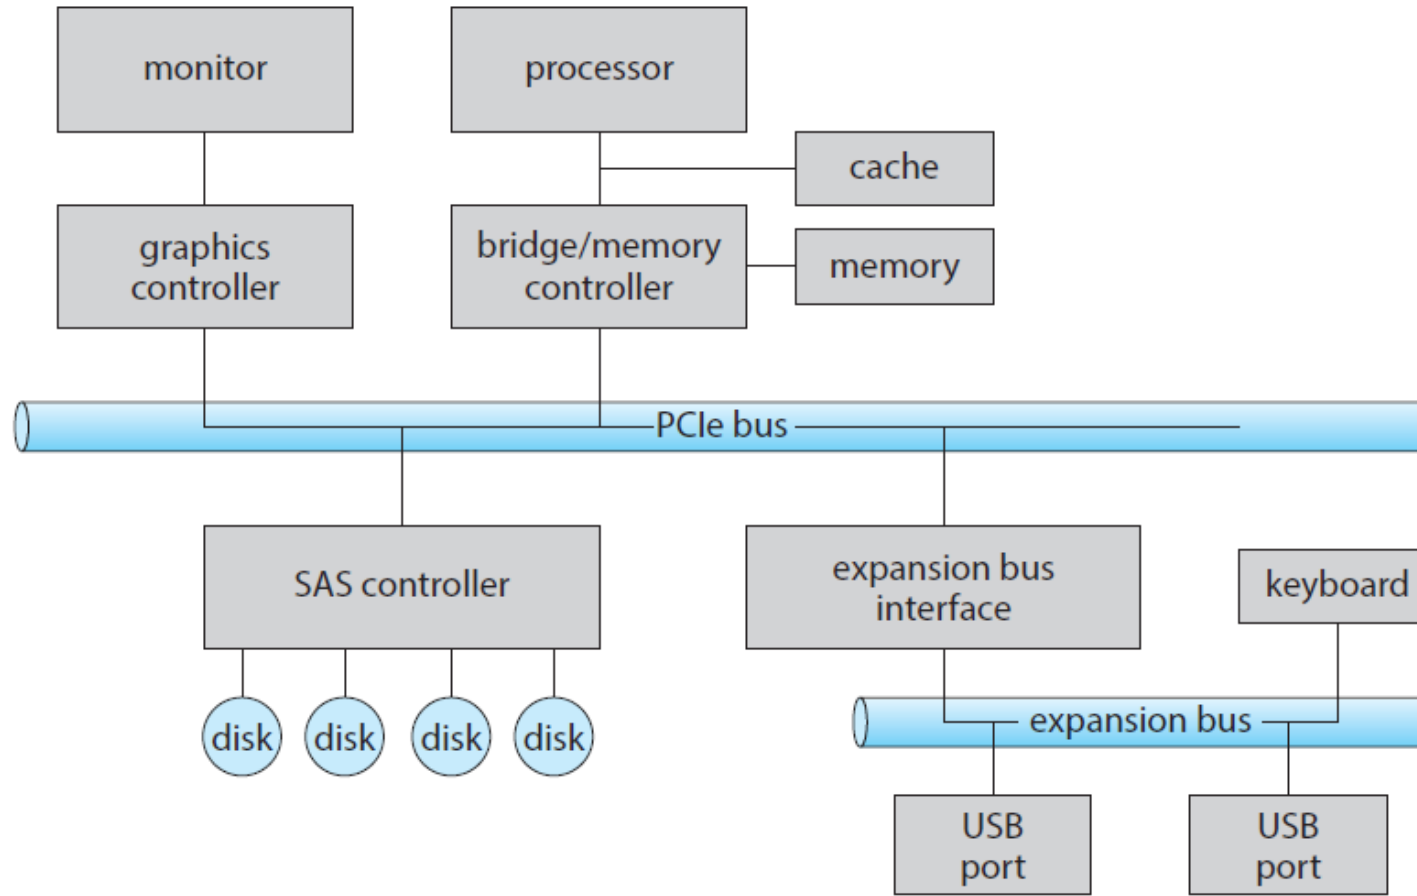


Figure 12.1 *A typical PC bus architecture.*



Ch. 12 I/O Systems

■ Memory-Mapped I/O

- How does the processor give commands and data to a controller to accomplish an I/O transfer?
- The controller has one or more registers for data and control signals.
 - *data-in* register: is read by the host to get input.
 - *data-out* register: is written by the host to send output.
 - *status* register: contains bits that can be read by the host.
 - *control* register: can be written by the host to start a command or to change the mode of a device.



Ch. 12 I/O Systems

■ Memory-mapped I/O

- If the device supports the memory-mapped I/O,
 - control registers are mapped into the address space of the processor.
- The CPU executes I/O requests
 - using the standard data-transfer instructions to read and write the device-control registers at their mapped locations in physical memory.

| I/O address range (hexadecimal) | device |
|---------------------------------|---------------------------|
| 000–00F | DMA controller |
| 020–021 | interrupt controller |
| 040–043 | timer |
| 200–20F | game controller |
| 2F8–2FF | serial port (secondary) |
| 320–32F | hard-disk controller |
| 378–37F | parallel port |
| 3D0–3DF | graphics controller |
| 3F0–3F7 | diskette-drive controller |
| 3F8–3FF | serial port (primary) |

Figure 12.2 Device I/O port locations on PCs (partial).



Ch. 12 I/O Systems

- Three types of I/O:
 - ***polling***: or *busy-waiting*:
 - reading the status register repeatedly until the busy bit becomes clear.
 - ***interrupt***:
 - CPU has a wire called the interrupt-request line.
 - If CPU detects an interrupt, it jumps to an ISR(*interrupt service routine*) to handle an interrupt.
 - The addresses of ISRs is specified in the *interrupt vector table*.
 - ***DMA: Direct Memory Access***:
 - used to avoid programmed I/O (one byte at a time).
 - useful for handling large data transfer.



Ch. 12 I/O Systems

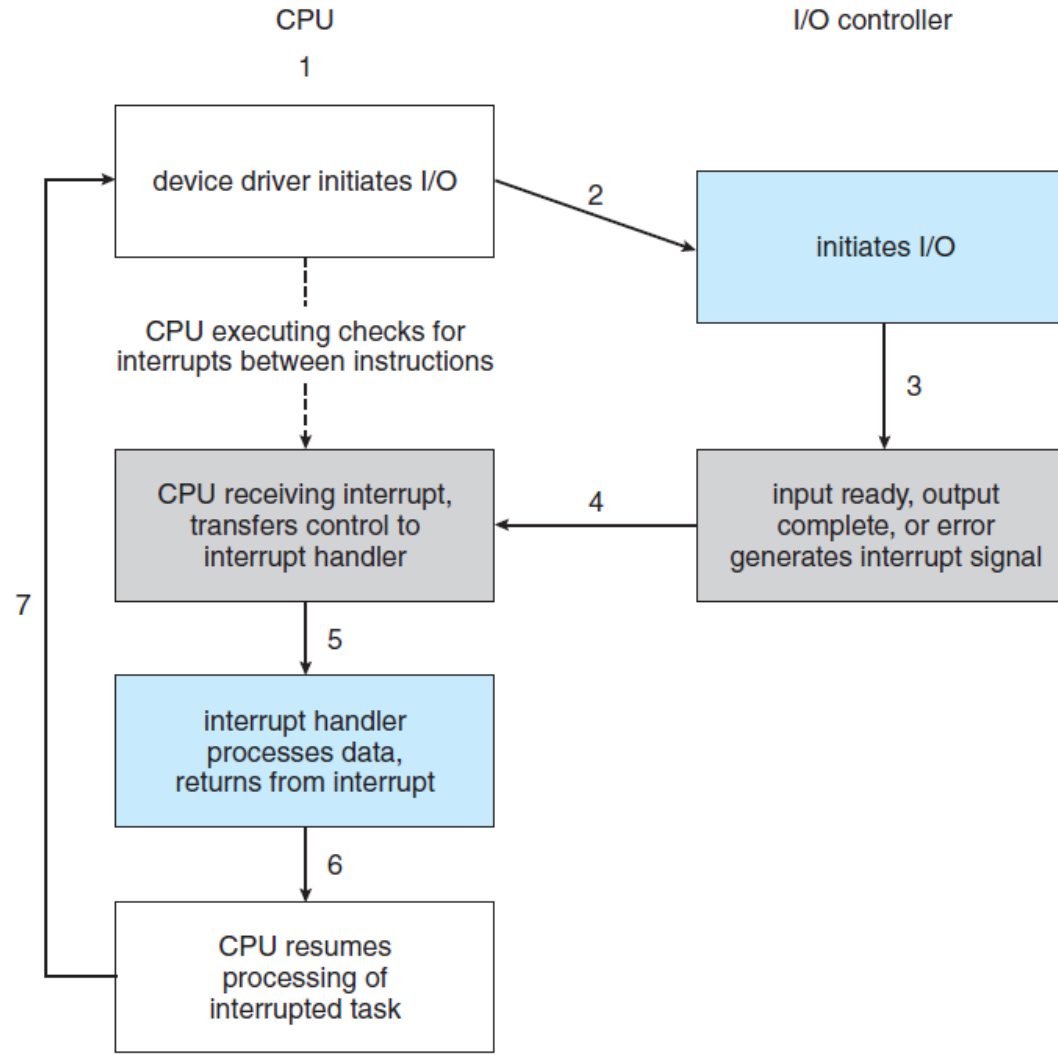


Figure 12.3 *Interrupt-driven I/O cycle.*



Ch. 12 I/O Systems

| vector number | description |
|---------------|--|
| 0 | divide error |
| 1 | debug exception |
| 2 | null interrupt |
| 3 | breakpoint |
| 4 | INTO-detected overflow |
| 5 | bound range exception |
| 6 | invalid opcode |
| 7 | device not available |
| 8 | double fault |
| 9 | coprocessor segment overrun (reserved) |
| 10 | invalid task state segment |
| 11 | segment not present |
| 12 | stack fault |
| 13 | general protection |
| 14 | page fault |
| 15 | (Intel reserved, do not use) |
| 16 | floating-point error |
| 17 | alignment check |
| 18 | machine check |
| 19–31 | (Intel reserved, do not use) |
| 32–255 | maskable interrupts |

Figure 12.5 *Intel Pentium processor event-vector table.*



Ch. 12 I/O Systems

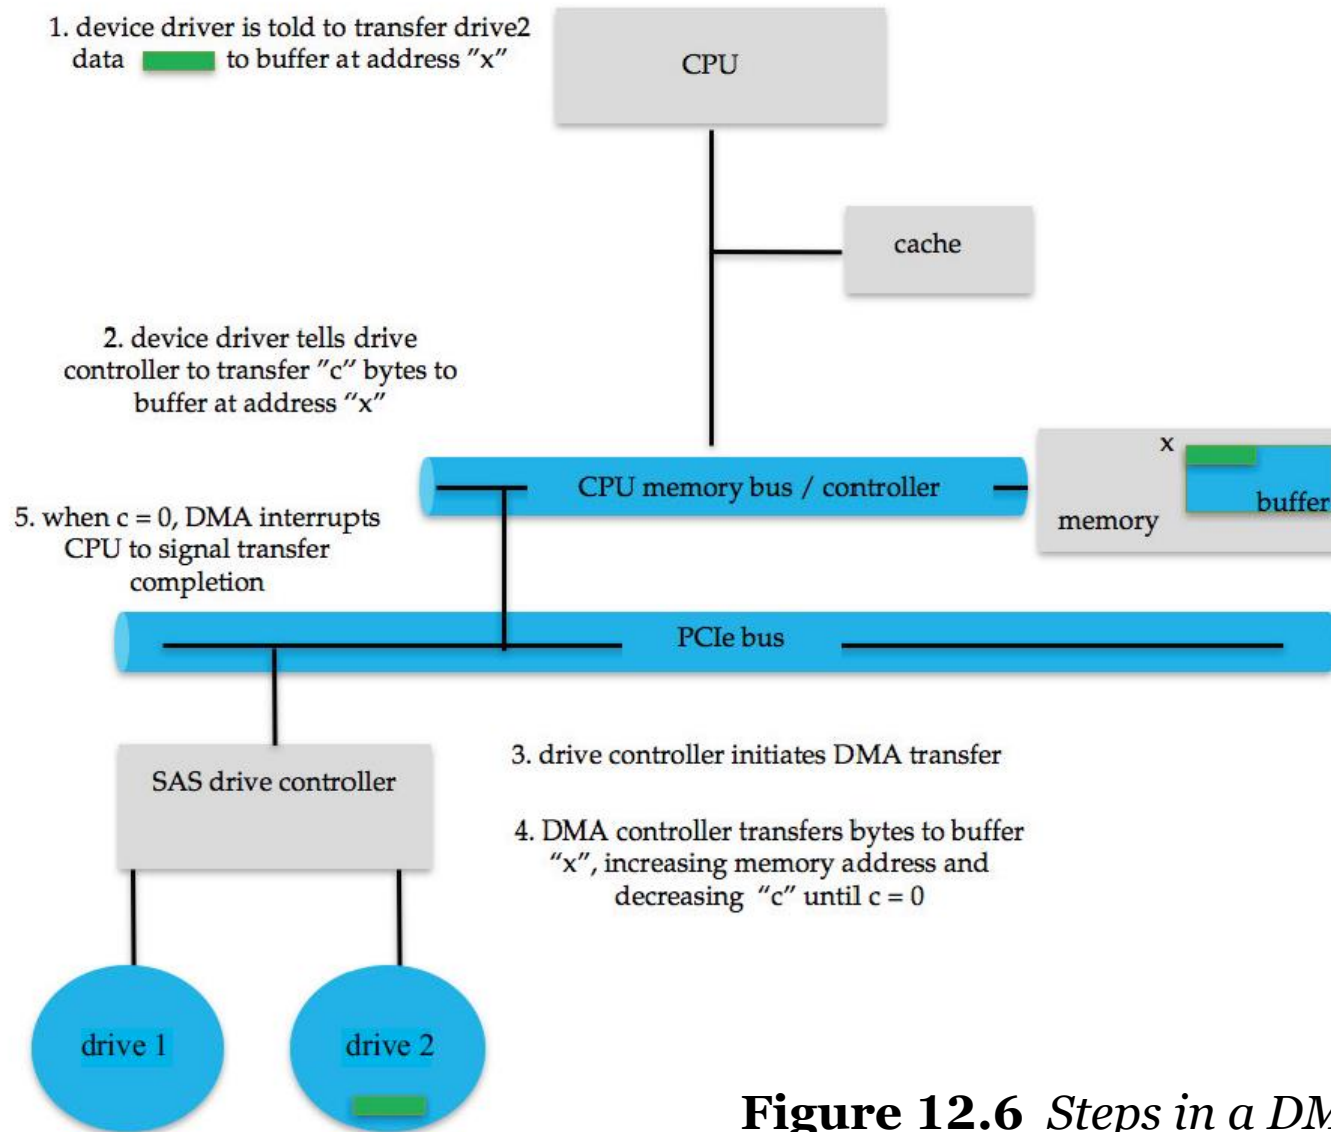


Figure 12.6 Steps in a DMA transfer.



Ch. 12 I/O Systems

- Blocking I/O .vs. Non-blocking I/O
 - *Blocking* I/O: a thread is suspended.
 - moved from *running* queue to *waiting* queue.
 - *Non-blocking* I/O: does not halt the execution of the thread.
 - e.g., receiving keyboard or mouse input in word processor.
 - returns as much as available.
 - *Asynchronous system call*: the thread continues to execute its code.



Ch. 12 I/O Systems

- The diff. between *non-blocking* and *asynchronous* system call:
 - a *non-blocking* read() call
 - *returns immediately* with whatever data are available,
 - the full number of bytes requested, fewer, or none at all.
 - An *asynchronous* read() call
 - *requests a transfer* that will be performed in its entirety
 - but will complete at some future time.



Ch. 12 I/O Systems

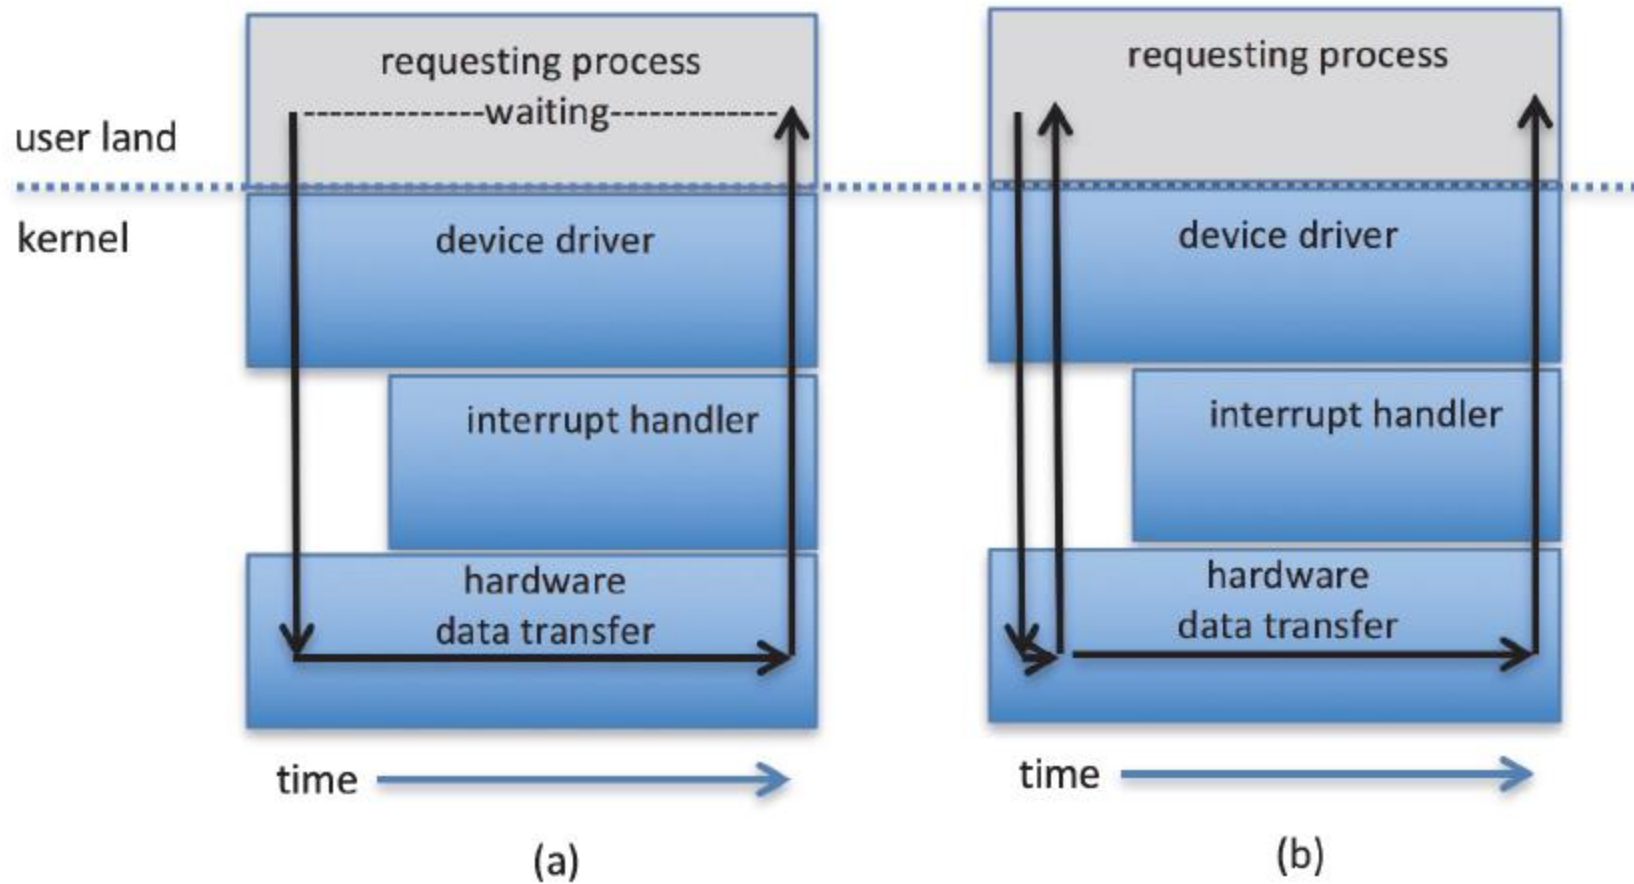


Figure 12.9 Two I/O methods: (a) synchronous (b) asynchronous.



Ch.13 File-System Interface

■ File System:

- provides the mechanism for
 - on-line storage of and access to both *data* and *programs*
 - of the *operating system* and *all the users* of the computer system.
- consists of two distinct parts
 - a collection of *files*, each storing related data
 - a *directory* structure, which organizes all the files in the system.



Ch.13 File-System Interface

■ Access Methods:

- *sequential access*:
 - Information in the file is processed in order,
 - one record after the other.
- *direct access*: relative access:
 - A file is made up of fixed-length logical records that
 - allow programs to read and write records rapidly in no particular order.



Ch.13 File-System Interface




Figure 13.4 *Sequential-access file.*

| sequential access | implementation for direct access |
|-------------------|---|
| reset | <code>cp = 0;</code> |
| read_next | <code>read cp;</code> <code>cp = cp + 1;</code> |
| write_next | <code>write cp;</code> <code>cp = cp + 1;</code> |

Figure 13.5 *Simulation of sequential access on a directed-access file.*



Ch.13 File-System Interface

- Directory Structure:
 - The directory can be viewed as
 - a symbol table that translates file names into their file control blocks.
 - The ways of organizing the directory structure:
 - *Single-Level Directory*
 - *Two-Level Directory*
 - *Tree-Structured Directories*
 - *Acyclic-Graph Directories*



Ch.13 File-System Interface




Figure 13.7 *Single-level directory.*

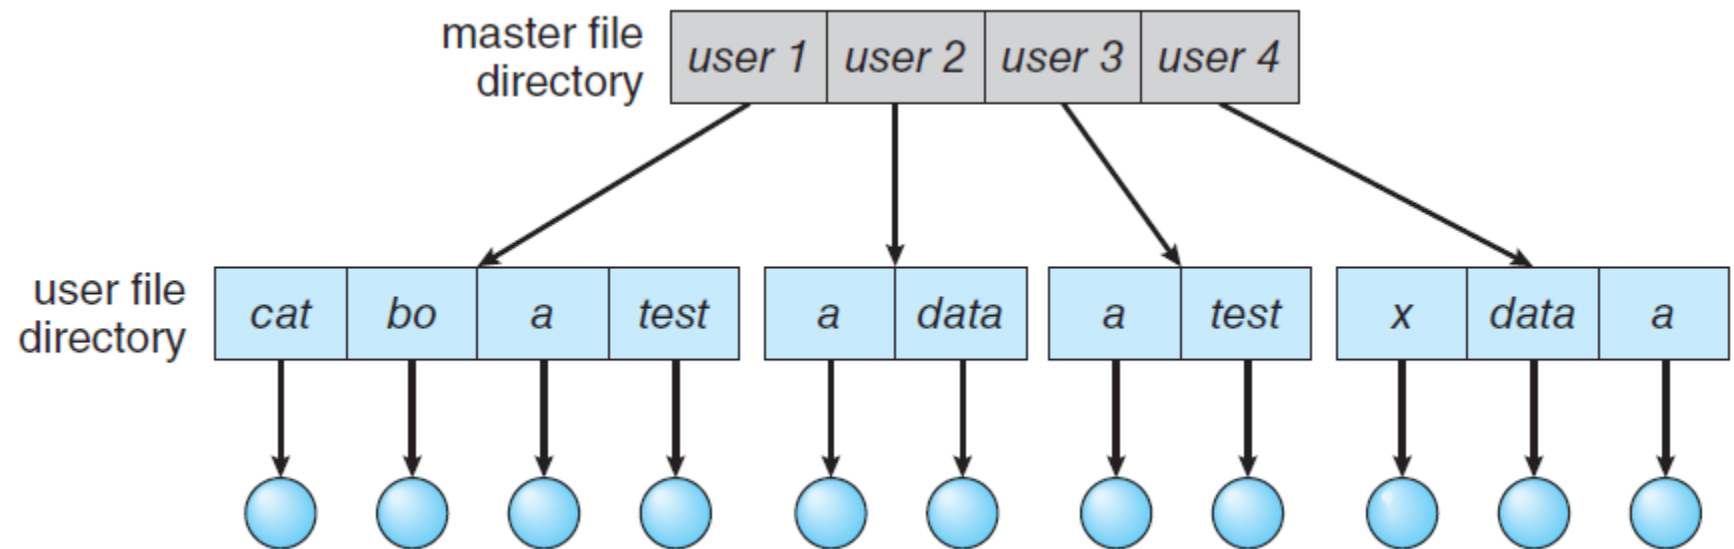


Figure 13.8 *Two-level directory.*



Ch.13 File-System Interface

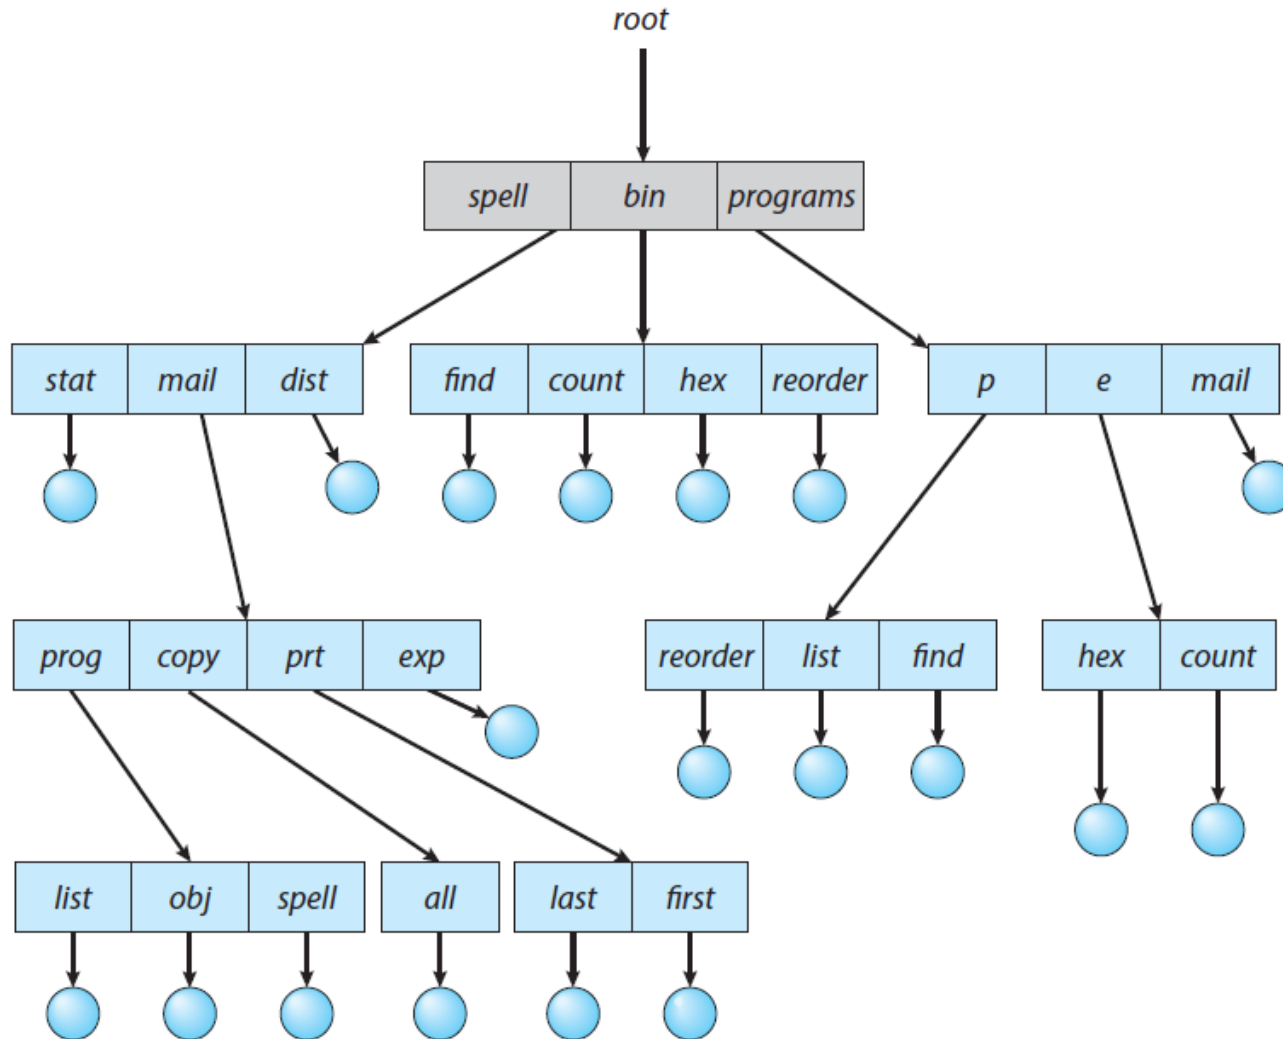


Figure 13.8 Tree-structured directory structure.



Ch.13 File-System Interface




Figure 13.10 *Acyclic-graph directory structure.*




Figure 13.11 *General graph directory.*

Ch. 14 File-System Implementation

- The file system itself
 - is generally composed of many different levels.

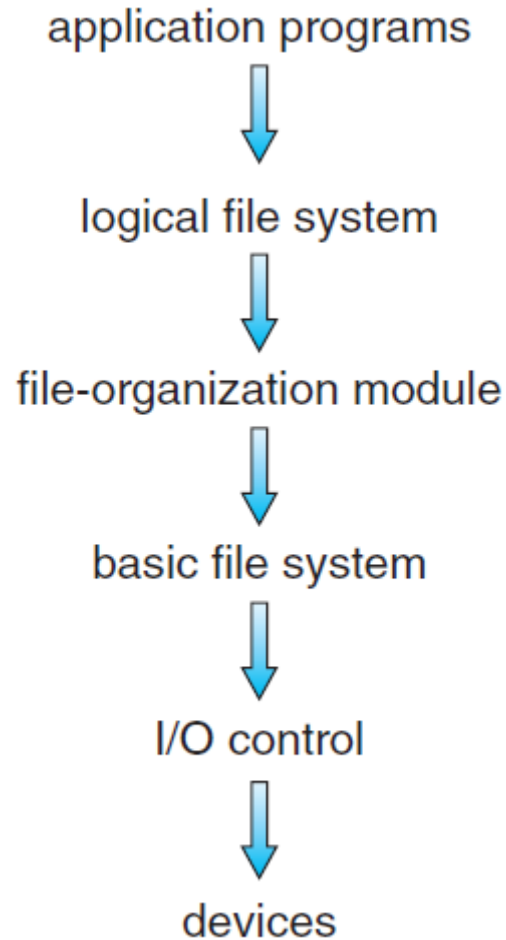


Figure 14.1 *Layered file system.*

Ch. 14 File-System Implementation

■ *Allocation Method:*

- The main problem of implementing the file system:
 - how to *allocate space to files*
 - so that storage space is *utilized effectively*
 - and files can be *accessed quickly*.
- Three major methods in wide use:
 - *Contiguous Allocation*
 - *Linked Allocation*
 - *Indexed Allocation*

- *Contiguous Allocation:*

- requires that each file occupy a set of contiguous blocks on the device.

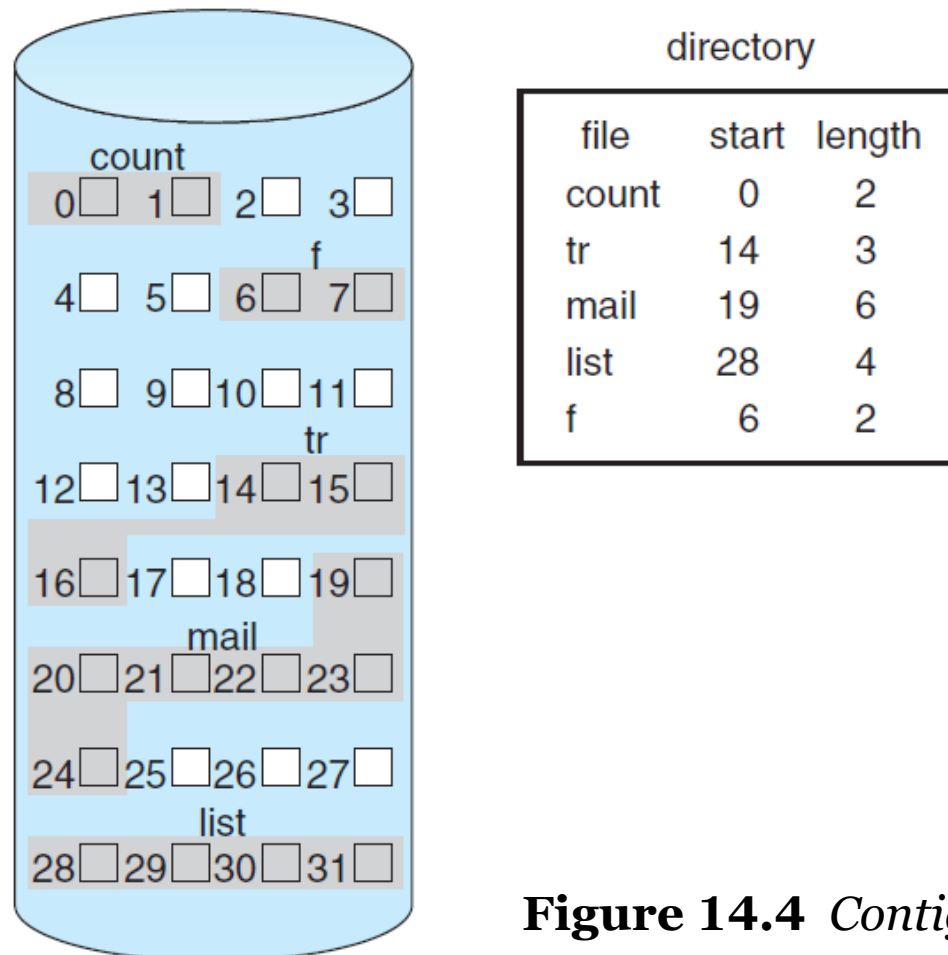


Figure 14.4 *Contiguous allocation of disk space.*

Ch. 14 File-System Implementation

■ *Linked Allocation:*

- The problems of contiguous allocation:
 - *external fragmentation.*
 - *need for compaction.*
- Linked Allocation
 - solves all problems of contiguous allocation
 - each file is a linked list of storage blocks
 - the blocks may be scattered anywhere on the device.



Ch. 14 File-System Implementation




Figure 14.5 *Linked allocation of disk space.*

Ch. 14 File-System Implementation

■ *FAT: File Allocation Table*

- The disadvantages of the linked allocation:
 - It can be used effectively only for sequential-access files
 - To find i -th block of a file, we must start at the beginning of that file.
- FAT: an important variation on the linked allocation
 - simple but efficient method using a *file allocation table*.
- A section of storage at the beginning of each volume
 - is set aside to contain the table.
 - the table has one entry for each block and is indexed by block number.



Ch. 14 File-System Implementation

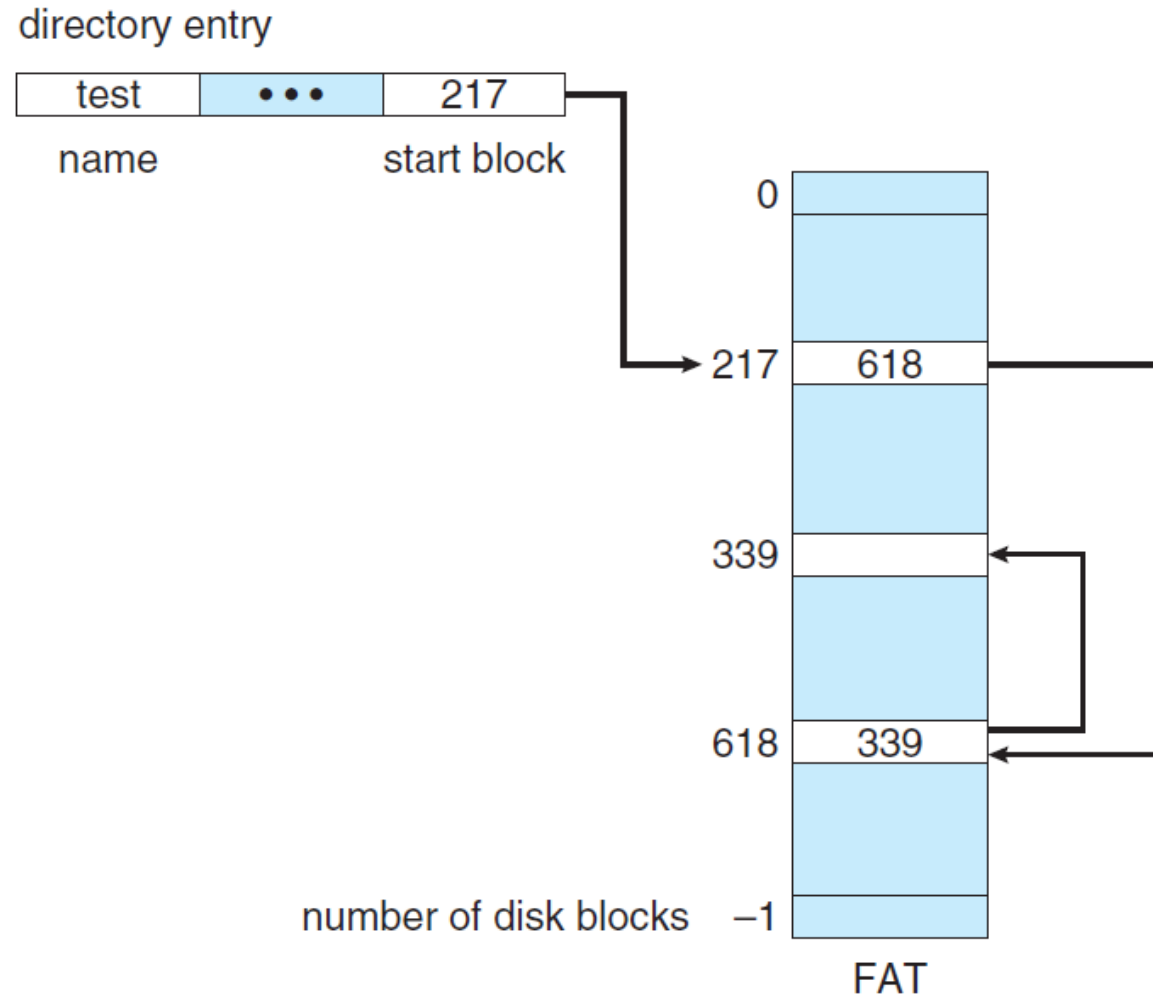


Figure 14.6 *File Allocation Table.*

Ch. 14 File-System Implementation

■ *Indexed Allocation:*

- The problems of linked allocation (without the FAT)
 - the pointers to the blocks are scattered with blocks
- Indexed Allocation solves this problem
 - by bringing all the pointers together into the *index block*.
- Each file has its own index block,
 - which is an array of storage-block addresses.
 - the i -th entry in the index block points to the i -th block of the file.



Ch. 14 File-System Implementation

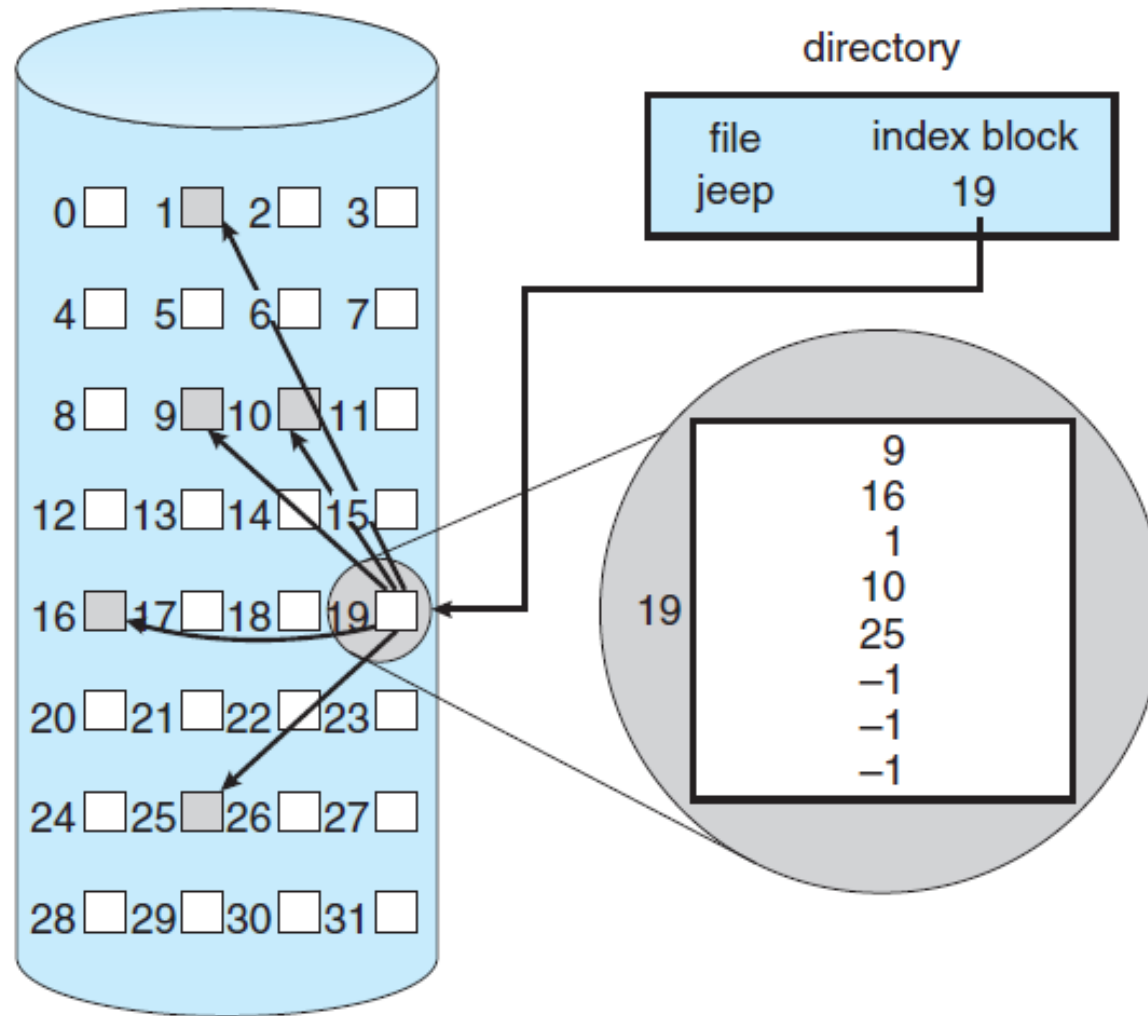


Figure 14.7 Indexed allocation of disk space.

Ch. 14 File-System Implementation

- Free-Space Management
 - To keep track of free disk space,
 - the system maintains a *free-space list*.




Figure 14.9 *Linked free-space list on disk.*

Any Questions?

